

Ch. 2 (contd.) **Applications and Layered Architectures**

Ch. 8 **TCP/IP**

TCP

- Segment format
- Ports and multiplexing/demultiplexing
- Connection management
- Window management; flow and congestion control
- Timer management
- Berkeley Sockets

Transport Protocols

Three important protocols:

1. UDP unreliable, connection-less

2. TCP reliable, connection-oriented we will focus on this

3. RPCs

Some Message Transport Requirements

If required, a **transport** protocol must be available to

- Guarantee message delivery
- Deliver messages in-order
- Identify and eliminate message duplicates
- Deliver messages regardless of size
- Provide sender-receiver sync
- Implement flow control (recvr → sender)
- Support multiple applications simultaneously on host

A transport protocol ***need not do all*** of those things

We merely ask that, if required, a transport protocol be available to do at least some of those things, and that ..

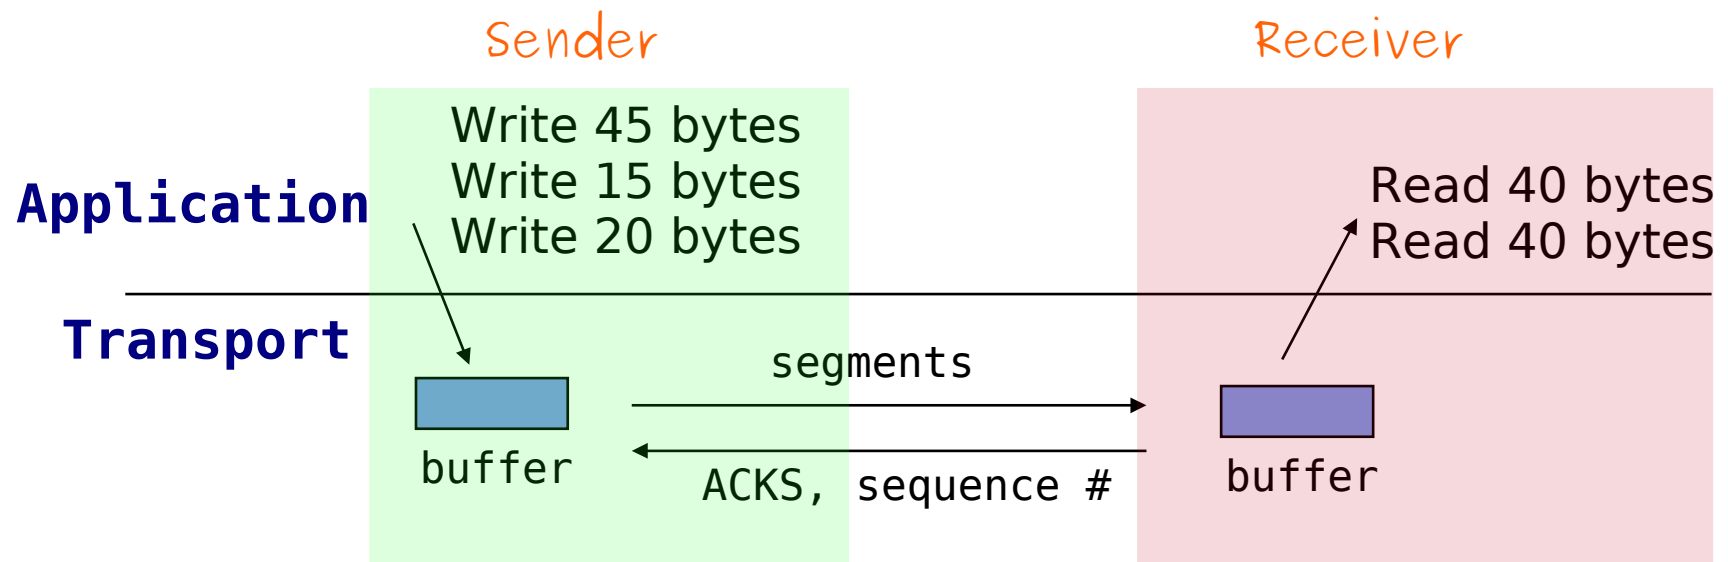
A suite of transport protocols ***cover all*** of those things *in the aggregate*

The Transmission Control Protocol (TCP)

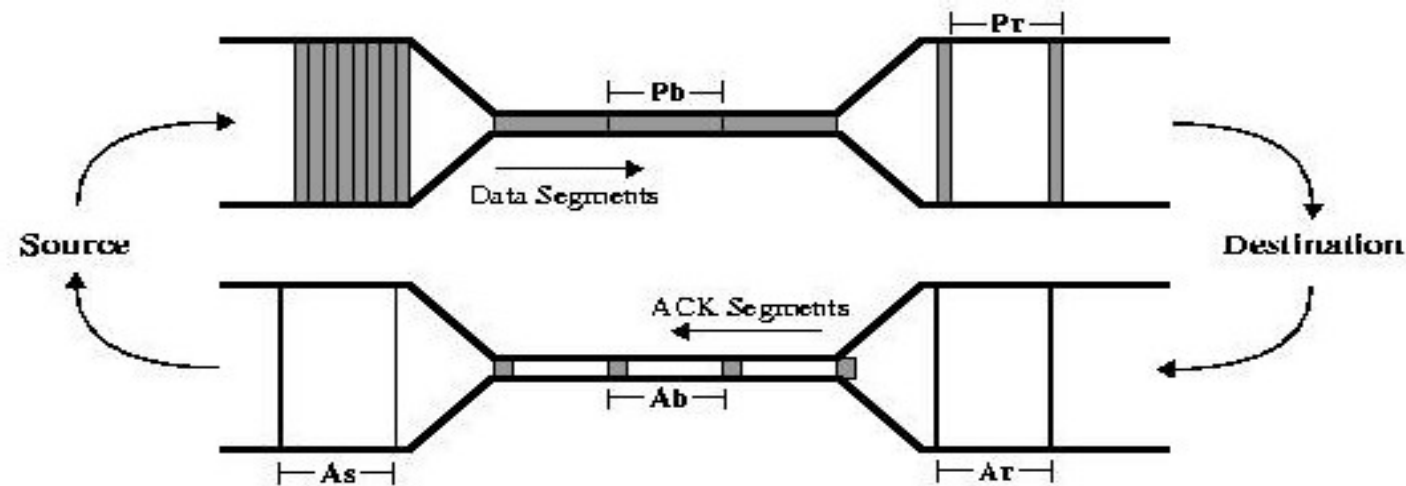
- TCP is *connection-oriented* transport layer protocol wherein
 - Sender & Receiver
 - first establish connection
 - set initial sequence numbers, sliding windows, timers
 - then, transfer data; use error recovery (by re-txing, if reqd.)
 - finally, tear down connection when done
 - You get *full-duplex, reliable byte stream service* between two processes in two computers *anywhere*
 - Sequence numbers track transmitted & received bytes
 - Error detection and retransmission allows recovery from transmission/reception errors
 - Sliding windows with feedback, used for flow & congestion control

TCP: Reliable Byte-Stream Service


- Stream Data Transfer
 - transfers contiguous stream of bytes across network, with no indication of boundaries
 - groups bytes into segments (**segments = packets**)
 - transmits segments as convenient (Push, if needed)
- Reliability
 - error control mechanism



TCP: Congestion & Flow Control



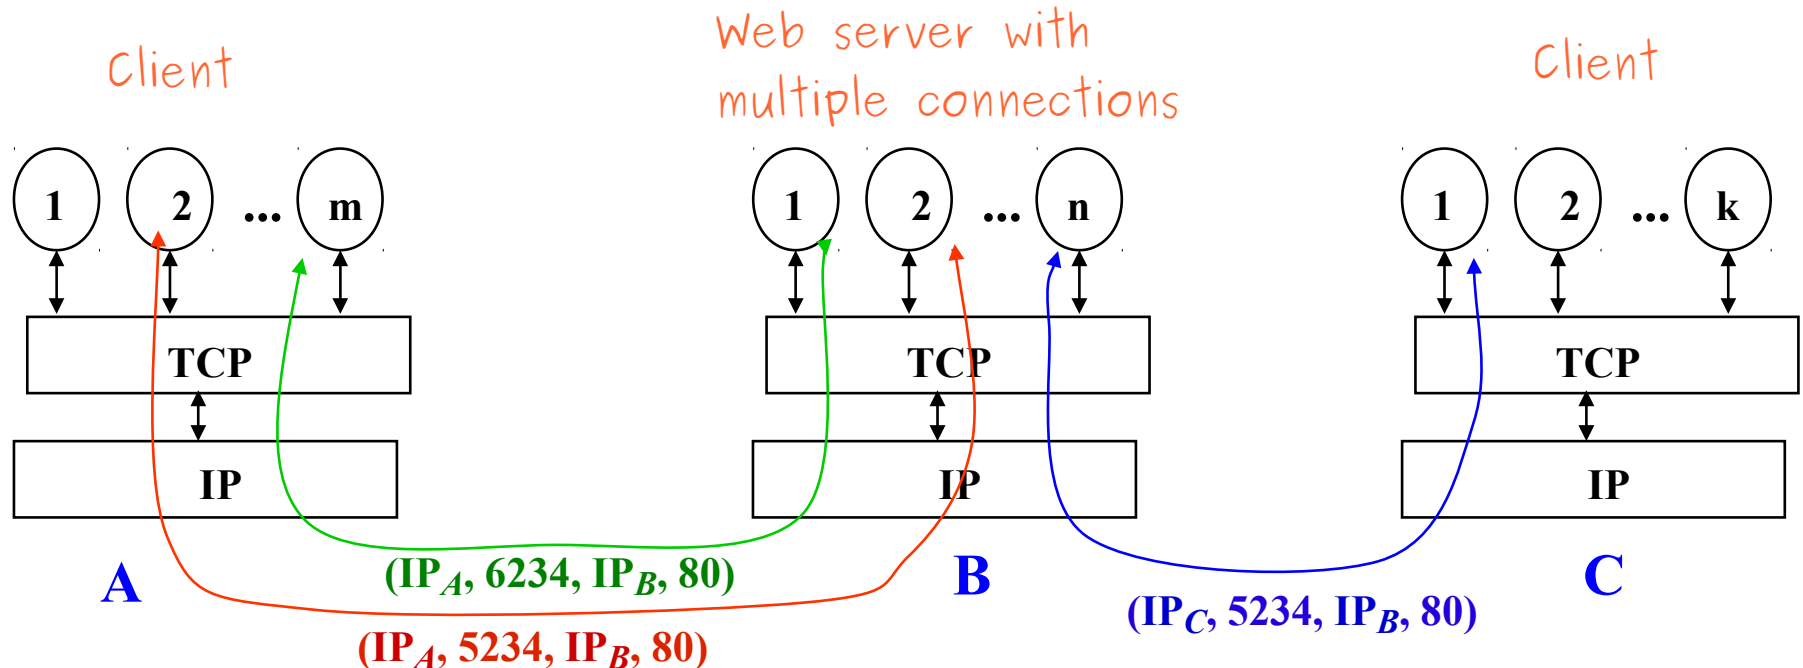
(a) Flow determined by Network Congestion



(b) Flow determined by Destination System

TCP: Multiplexing / Demultiplexing Connections

- TCP connection specified uniquely by
 - **(src IP addr, src port #, dest IP addr, dest port #)**
 - *above 4-tuple is called TCP's demultiplexing key*
 - *Caution! connection may have more than 1 incarnation*
- Multiplexing allows multiple end-to-end apps. simultaneously
- Arriving segment handed to appropriate app. based on connection 4-tuple



TCP: Ports & Port Numbers

- TCP Ports are **virtual** (not physical)
- Router port = router's physical interface
- Server processes (apps) *listen* on ports,
- Client processes (apps) *send* messages to ports

Question: How does client know server's port number?

Answer: Server processes (daemons) use *well-known ports* (i.e. port numbers less than 1024)

Examples:

<u>App Process</u>	<u>Port #</u>
TCP telnet	23
TCP ssh	22
TCP ftp	21
TCP http	80
UDP DNS	53
UDP talk	517

For a more complete list, see:
<http://www.networksorcery.com/enp/protocol/ip/ports00000.htm>

etc.

TCP: Ports & Port Numbers (contd.)

Question: What if the ports are not so *well known*?

Answer: Sometimes, server process uses non-std port, or uses different port number on each start up. Then, it registers its port # with a **port-mapper** service. The port-mapper service listens on a **well-known port** (UDP or TCP port number **111**)

- client contacts port-mapper to get server's port #
- port-mapper tells client server's port #

Question: How does server know client's port number?

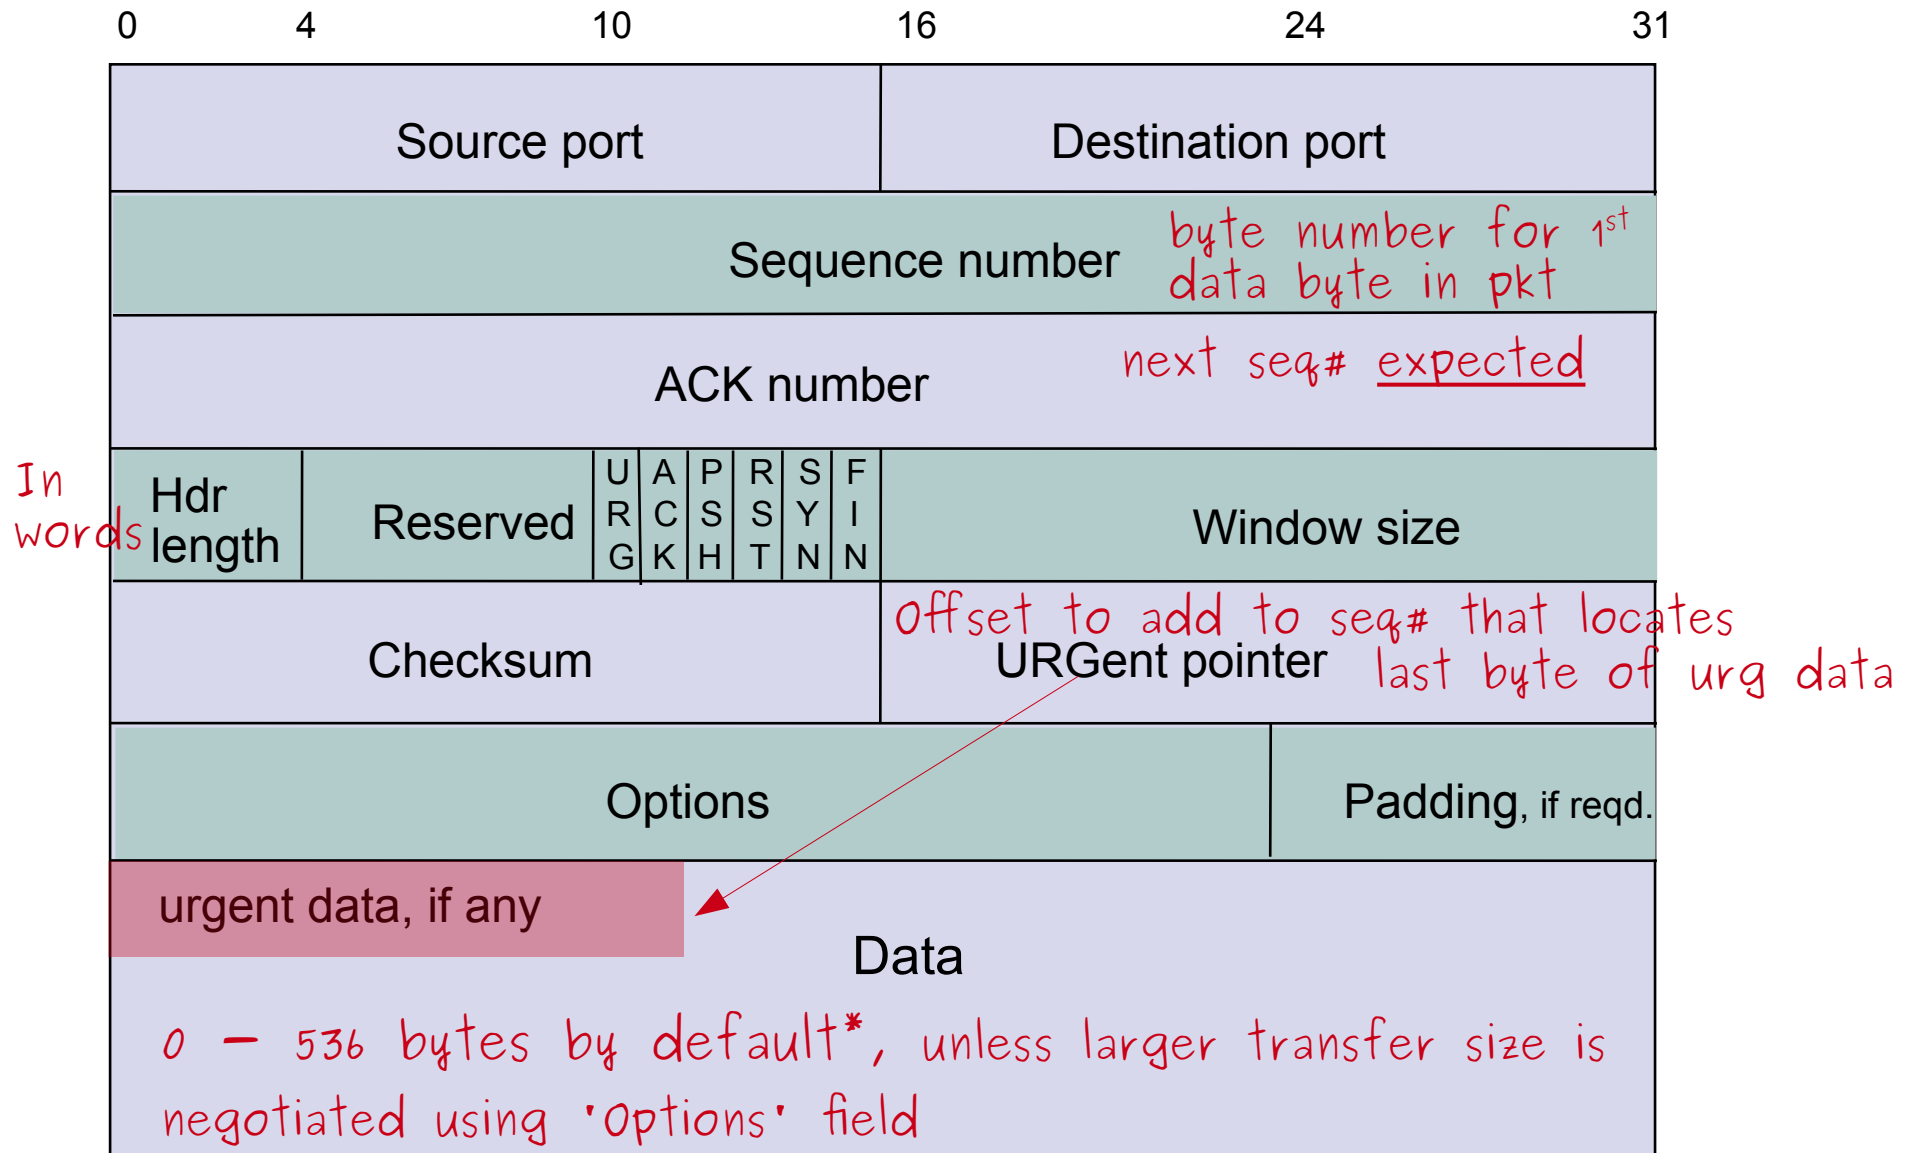
Answer: Client's port number is in client's service request (UDP or TCP packet header) sent to server

- **Port Implementation:**

Incoming packets are put in finite-buffer FIFO queues

- pkt drops are possible
- App process blocks if it tries to read empty queue

TCP: Segment Format (Header & Data)



*IP layers are reqd to carry at least 576 bytes of IP packet **without** fragmentation

TCP Header Fields Explained

(refer to figure on previous page)

Port Numbers

- A **socket** identifies a connection endpoint
 - IP address + port #
- Connection is specified by a *socket pair*

Acknowledgement Number

- SN of next byte expected by receiver
- Acknowledges that all prior bytes in stream have been received correctly
- Valid if ACK flag is set

Sequence Number

- Byte count
- First byte in segment
- 32 bits long
- $0 \leq SN \leq 2^{32}-1$
- Initial sequence number is selected during connection setup

Hdr Length

- 4 bits long
- Length of header in multiples of 32-bit words
- Minimum hdr length = 20 bytes, max = 60 bytes

TCP Header Fields Explained (contd.)

Control

- 6 flags, 1 bit per flag
- **URG**: urgent pointer flag
 - Urgent message end
= SN + **urgent pointer**
- **ACK**: ACK flag
- **PSH**: override TCP buffering
- **RST**: reset connection
 - Upon receipt of RST, connection is terminated and application layer notified
- **SYN**: establish connection
- **FIN**: close connection

Window Size

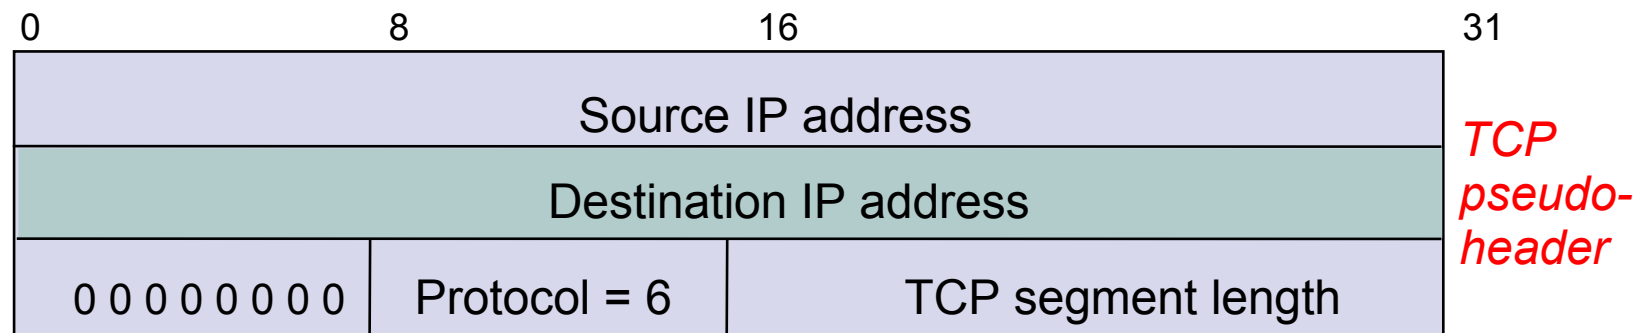
- 16 bits to advertise window size
- Used for **flow control**
- Sender will accept bytes with SN from ACK to ACK + window
- Maximum window size is 65535 bytes

TCP Checksum *Home reading – look this up*

- Internet checksum method
- TCP pseudoheader + TCP segment

TCP Header Fields Explained (Contd.)

Checksum

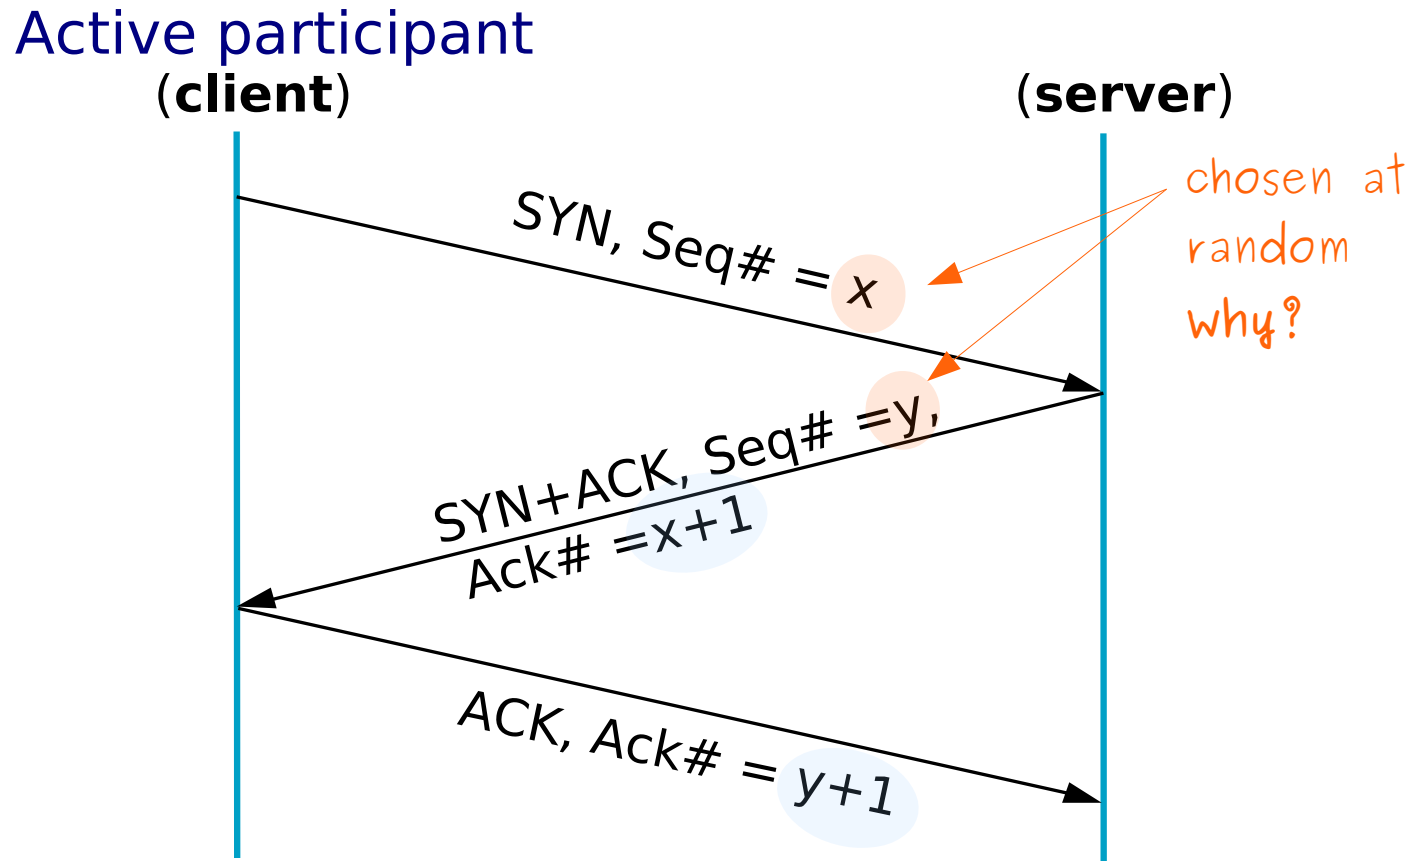


NOTE: *UDP and TCP use similar checksum computation. Read on your own.*

Options

- Variable length
- **NOP** (No Operation) option to pad TCP header to multiple of 32 bits; spacer betwn options
- **Time stamp** option for RTT measurements
- **Maximum Segment Size (MSS)** option specifies largest segment size acceptable to receiver
- **Window Scale** option increases TCP window size from 16 to 32 bits

TCP: Connection Establishment and Termination



A 3-way handshake for connection establishment

TCP: Connection Establishment and Termination (contd.)

Two issues to consider:

*1.1 Why exchange seq #s at connection establishment?
(i.e. Why not simply start from seq# 0 on each new connection?)*

1.2. Why not tie initial sequence number to local timer?

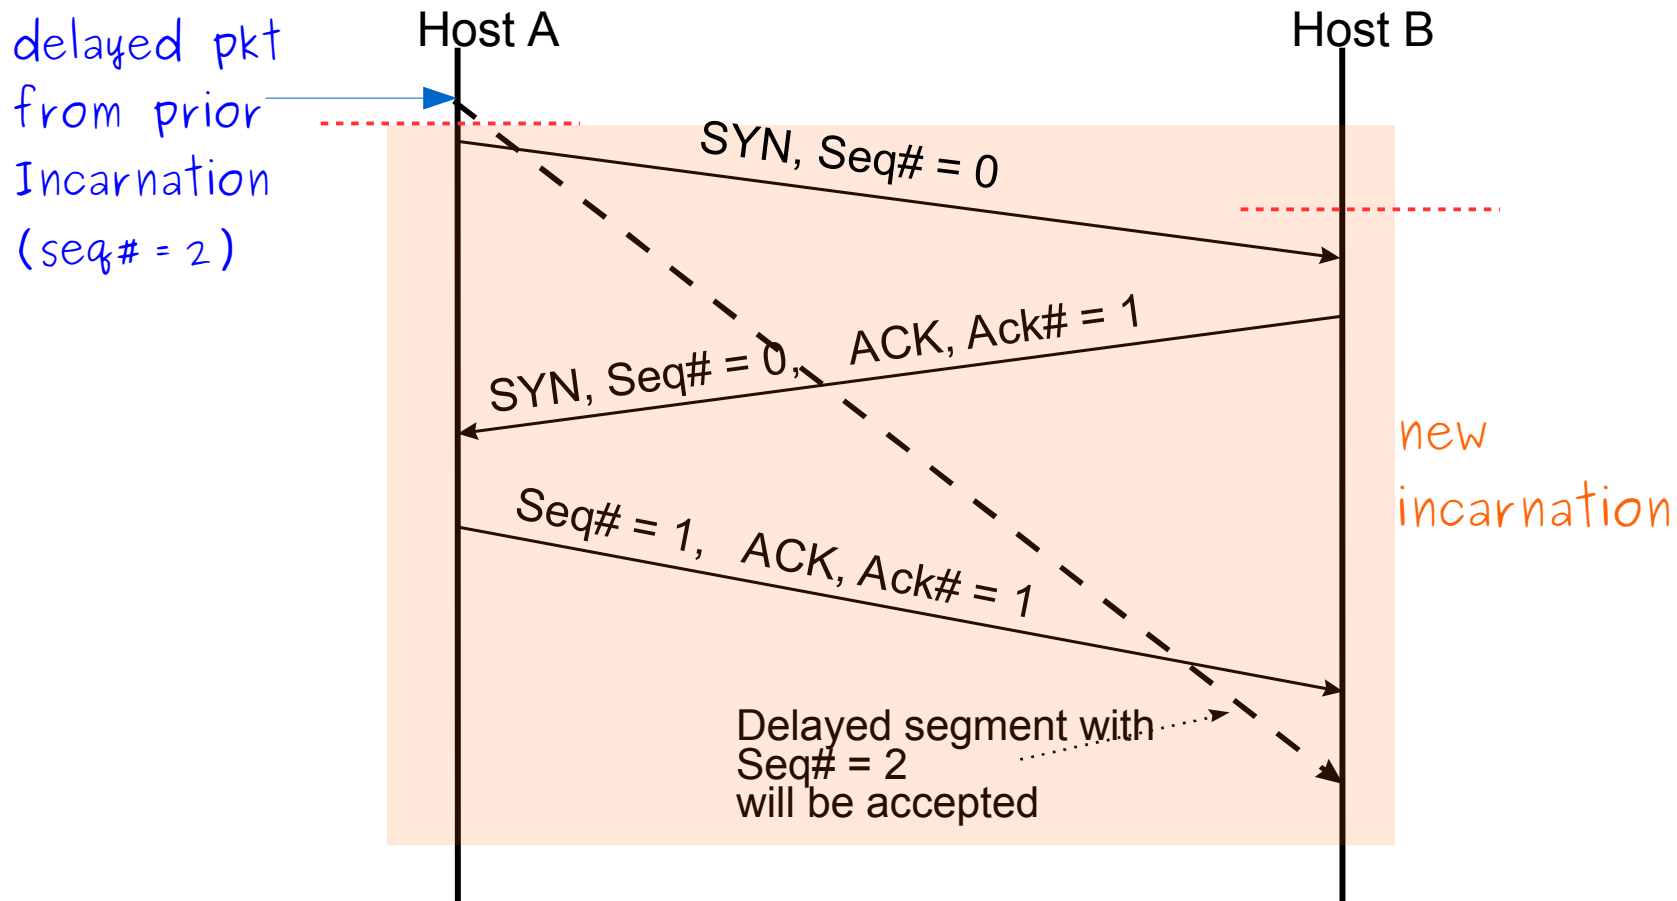
See: http://wiki.cas.mcmaster.ca/index.php/The_Mitnick_attack

*2. What about the 'two armies' problem?
(during connection tear-down phase)*

TCP: Connection Establishment and Termination (contd.)

1.1 Why not simply start from seq# 0 on each new connection?

Because, this may happen:



TCP: Connection Establishment and Termination (contd.)

Sequence Numbers

- Select initial sequence numbers (ISN) to avoid overlap with sequence numbers of prior connections
- Use ~~local clock~~ to select ISN sequence number
- Time for sequence nos. to wrap around should be greater than the maximum lifetime of a segment (**MSL**);
Typically MSL=120 seconds


- High bandwidth connections pose a problem

- $2^n > \frac{2 * \text{MSL sec} * R \text{ bytes/sec}}{}$

total sequence nos.

Bytes transmitted in 2MSL period

TCP: Connection Termination



2. The 2 armies problem (“Computer Networks, 5ed”, A. Tanenbaum):
a connection termination conundrum

Review: Sliding Window Protocol with Finite Seq Numbers




Figure 1

Review: Sliding Window Protocol with Finite Seq Numbers

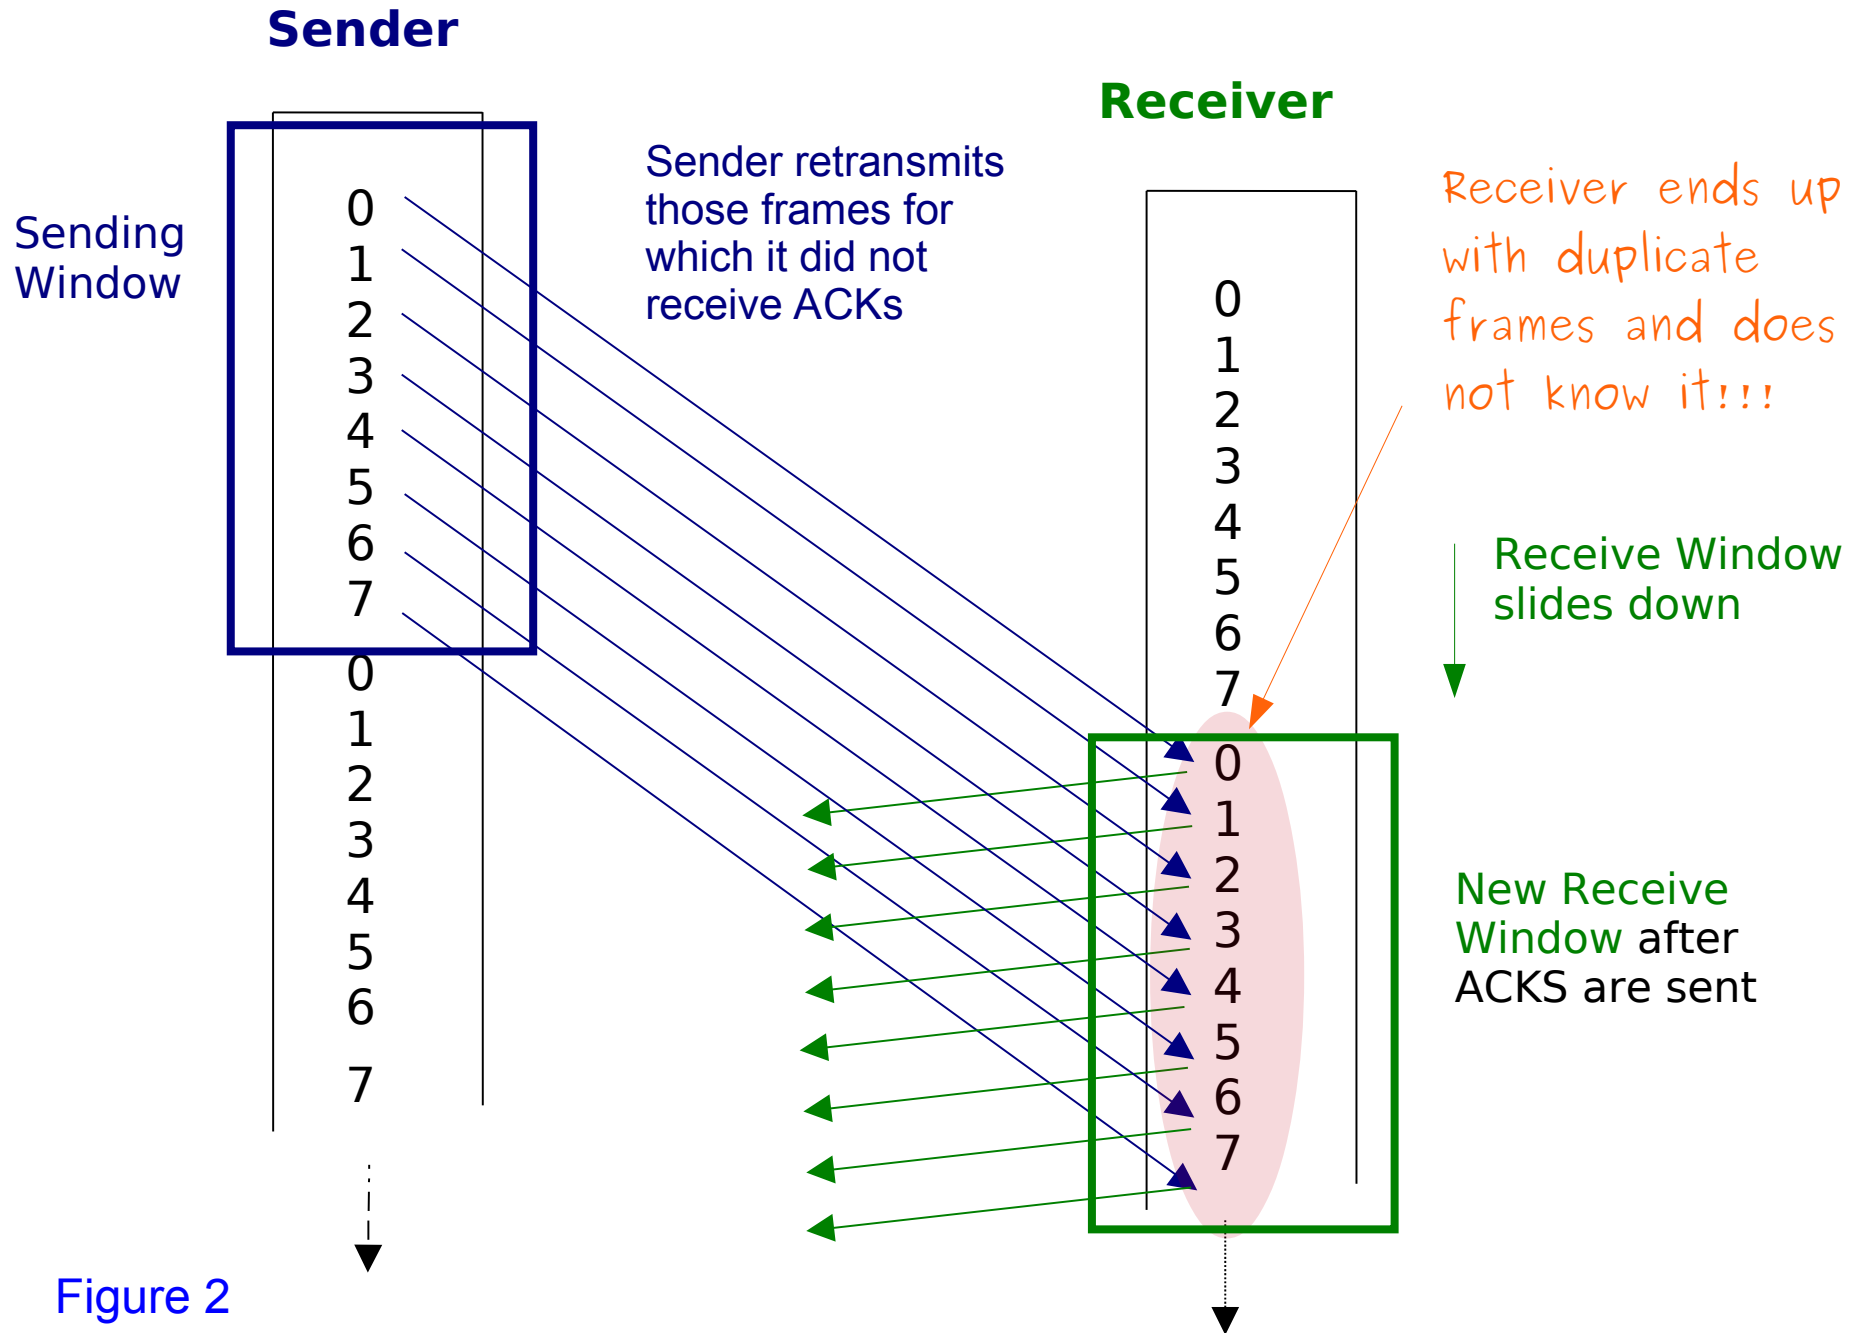


Figure 2

TCP: More On Sequence Numbers

The sliding window rule:

Receiver's successive windows must never overlap!

Let us see how TCP fares here ..

TCP's **Sequence Number** field \Rightarrow 32 bits

TCP's **Advertised Window** field \Rightarrow 16 bits

Therefore, no overlap possible between 2 successive advertised windows

*But is that enough to prevent **Sequence Number** wraparound problems that may arise?*

TCP: Sequence Number Wraparound

Consider from S's perspective..

1. S sends segment with **sequence number x** to R
2. S times out on **x** (no ACK recvd) .. **x** is delayed in transit, but *neither S nor R know that!*
3. S retransmits **x** to R

From R's perspective:

1. R recvs retransmitted **x** from S, sends ACK, then rolls down its window
2. R doesn't know that original copy of x is floating around on the net
3. Now, *some* time later seq nums at S roll over, S sends **new** (rolled over) **x** to R. Then, out of the blue ...
5. R receives very first copy of **x** sent by S, before it receives **new x**

TCP: Sequence Number Wraparound (contd.)

Bandwidth Time until Wraparound (of 32-bit sequence number)

Cable (10 Mbps) 57 minutes

Ethernet (100 Mbps) 6 minutes

STS-12 (622 Mbps) 55 seconds (!) .. *less than MSL*

STS-48 (2.4 Gbps) 14 seconds (!) .. *less than MSL*

TCP: Sequence Number Wraparound (contd.)

Solution?

Sending TCP process

- inserts 32-bit timestamp in **options** field of segment


Receiving TCP process

- checks timestamp as well as sequence number
- ➔ if segment ***n*** with timestamp smaller than segment ***n-i*** arrives, segment ***n*** is discarded

Why?

Because segment with lower valued timestamp is a delayed segment that just appeared at recvr

TCP: Advertised & Effective** Windows



Flow control using Advertised Sliding Windows

** definition of effective window will be updated later when we study congestion

TCP: Size of Advertised Window

Also, recall ...

- *To increase utilization and throughput, window size must be at least = ?*

Now, let us run the numbers ...

Advertised Window (**awnd**) field in TCP header is 16 bits
awnd is expressed as # of bytes

Therefore, **awnd** size is 2^{16} bytes (max)
= 64 KB (max)

Also, typically, RTT for cross-country segment = 100 ms

Now, refer to table on next slide

TCP: Size of Advertised Window (contd.)

Bandwidth	RTT Delay-Bandwidth Product (when RTT = 100ms)
Slow DSL (1.5 Mbps)	18 KB
Cable (10 Mbps)	122 KB (!) .. >> Max. Advertised Window
Ethernet (100 Mbps)	1.2 MB (!) .. >> Max. Advertised Window
STS-12 (622 Mbps)	7.4 MB (!) .. >> Max. Advertised Window
STS-24 (2.4 Gbps)	29.6 MB (!) .. >> Max. Advertised Window

So, for better line utilization, receiver must advertise a window larger than 64KB

How?

Use **options** field in TCP header to provide scaling (multiplication) factor for **awnd** size field

TCP: A Tiny Problem

One problem ...

- Suppose recvr's **AdvertisedWindow** (**awnd**) reaches 0
 - So, sender stops sending
 - Now, neither sender, nor receiver are sending anything to each other *deadlock?*
-

- Suppose now:
Recvr's application reads from receive buffer
(*and therefore frees up 'receive buffer' space*)

How does sender know new size of awnd ?

TCP: Solution to the Tiny Problem

Smart sender/dumb recvr rule...

When recvr's **awnd** reaches 0

- Sender sends segment with 1 byte of data once every **x secs**
- Receiver sends ACK segment in response

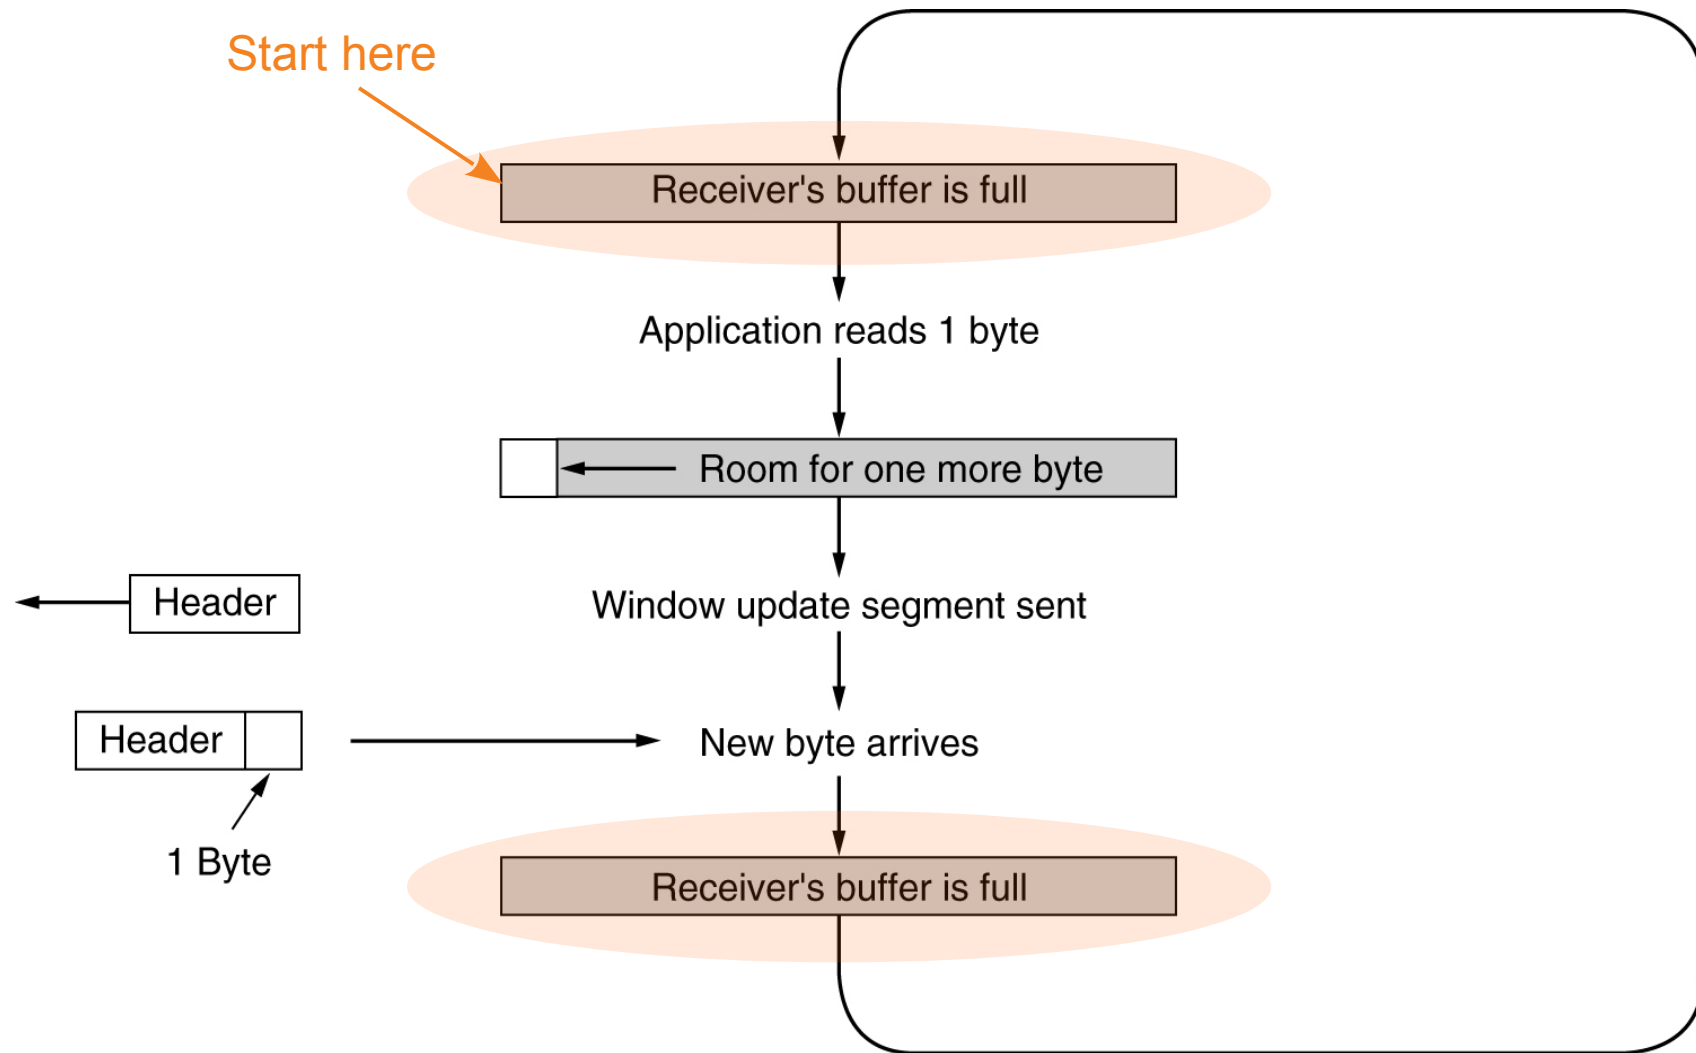
From **ACK#** field in TCP hdr of response segment,

- sender knows if pkt has been accepted/rejected

AND

- **awnd** field in TCP header of response segment tells sender the receiver's new advertised window

TCP: The *Silly Window Syndrome* (SWS)



Receiver's Window size stuck at 1 byte

TCP: The *Silly Window Syndrome* (contd.)

Two complementary approaches to SWS:

1. **Clark's solution** (on recvr's side)-

Receiver may send new **awnd** only if

- a. receive buffer has available space equal to MSS**
(negotiated at connection establishment time), **or**
- b. receive buffer is half empty

2. **Nagle's algorithm**

(= self-clocking solution on sender's side) .. *next slide*

******The Max Segment Size (MSS) is the maximum size of the TCP Payload (TCP, IP, MAC headers/trailers are excluded)

TCP: The *Silly Window Syndrome* (contd.)

2. Nagle's algorithm (= self-clocking solution):

when application produces data to send

if (**data-to-send** \geq **MSS** and **EffectiveWindow** \geq **MSS**)

then send a full segment

else

if there is unACKed data *in flight*

then buffer newly produced data until ACK arrives

else

send all new data now

Note:

1. Algorithm adjusts to RTT

- Short RTT - send frequently at low efficiency
- Long RTT - send less frequently at greater efficiency

2. If application cannot afford delay introduced by above algorithm, set **TCP_NODELAY** option when establishing socket .. this disables Nagle's algorithm altogether

TCP: The *Silly Window Syndrome* (contd.)

- Question to ask yourself:

*How does an application like **telnet** or **ssh** that generates data at a low rate, work with Nagle's algorithm?*

- Answer:

- *Trace through Nagle's algorithm to see how the 1st and the 2nd segments are transmitted*
- *Two successive 'small' writes to the TCP socket cause the 2nd TCP segment to wait while the ACK to the 1st segment is pending*
- *Causes delay of upto 500 ms!*
- *Better to use TCP_NODELAY option to disable Nagle's algorithm*

TCP: State Transition Diagram

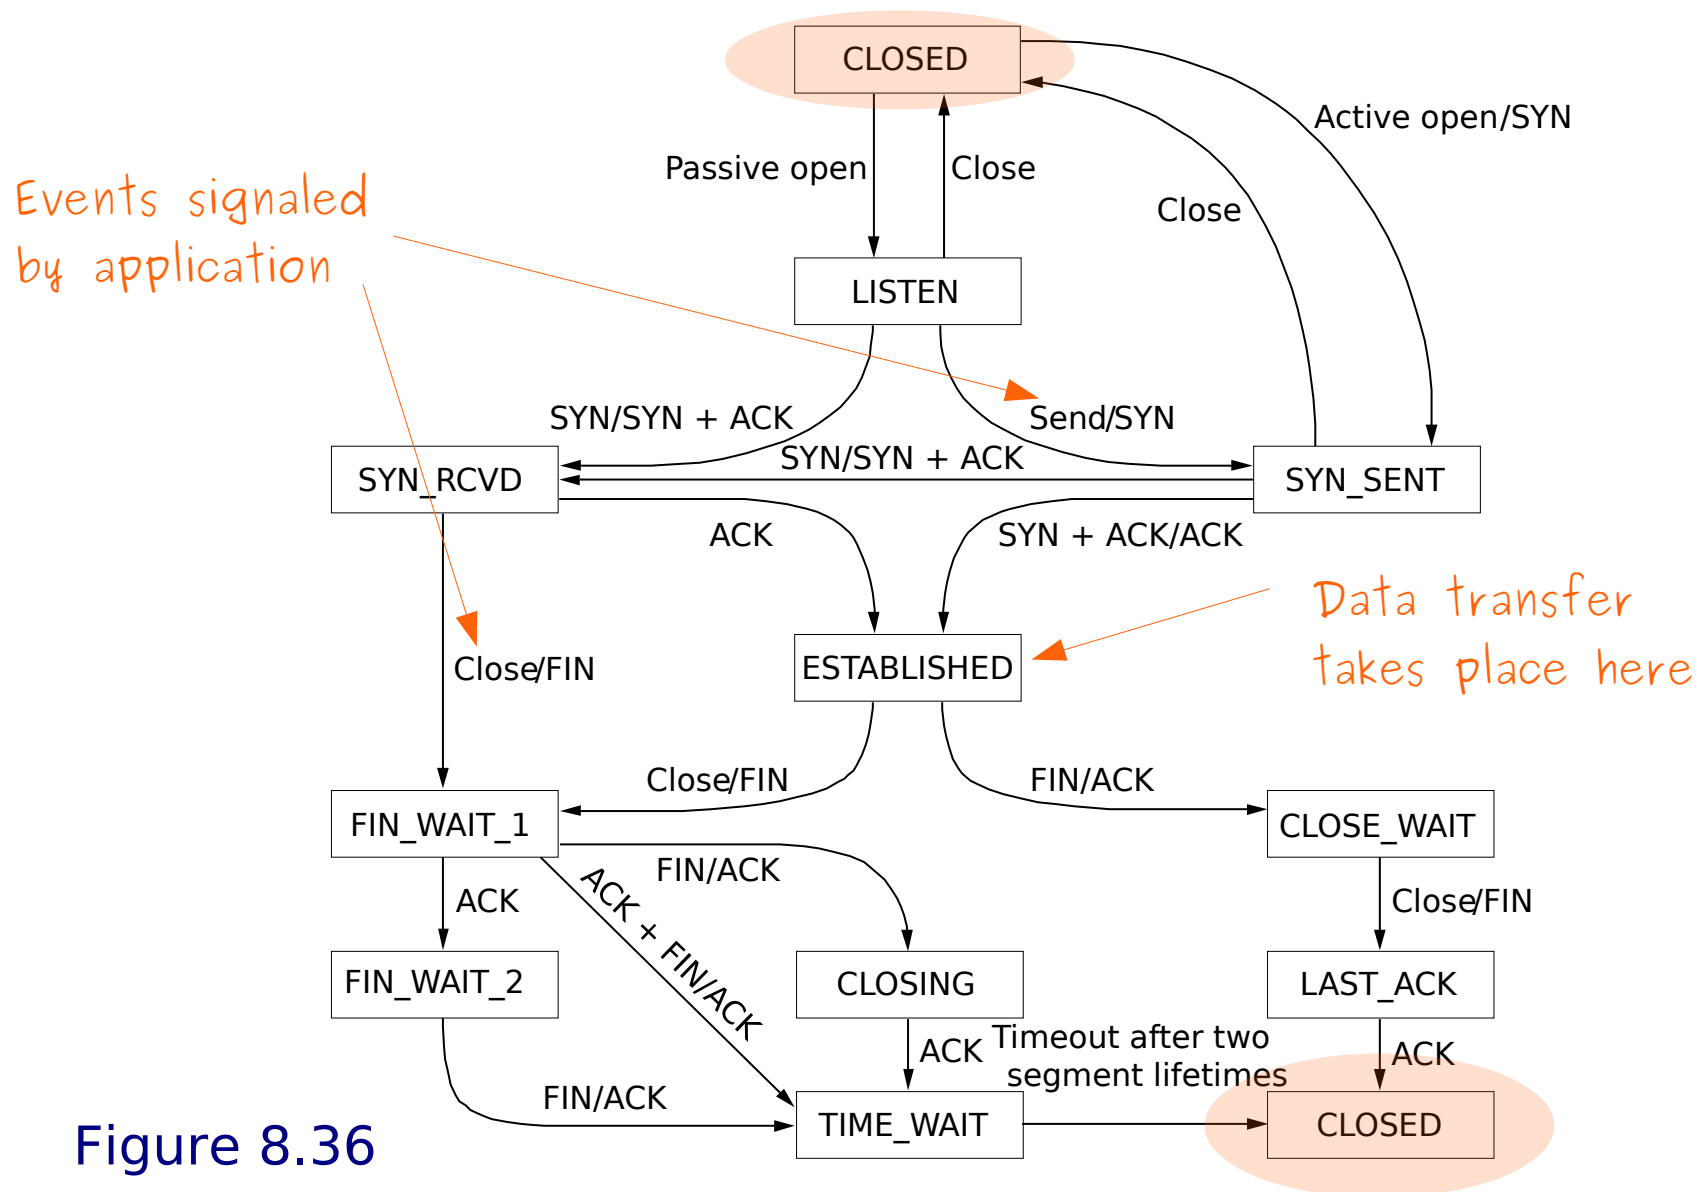
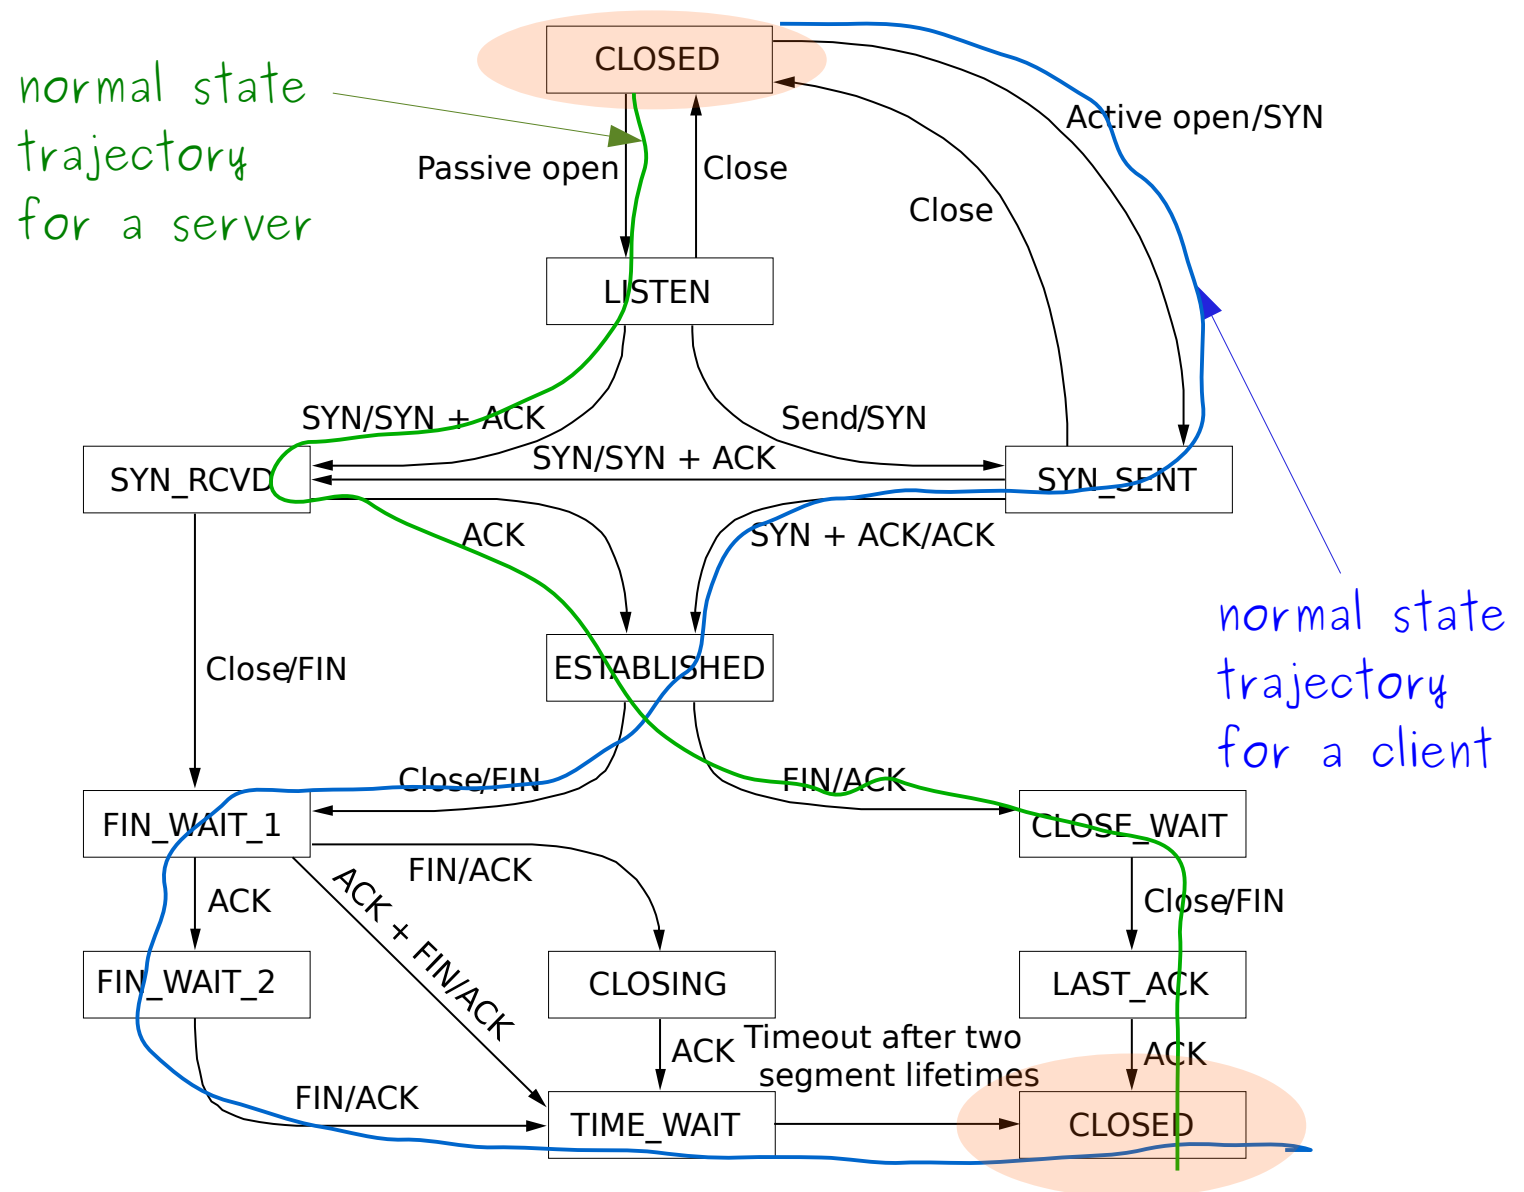


Figure 8.36

TCP: State Transition Diagram



TCP: **State Transition Diagram** (contd.)

Some issues to consider:

1. What if client's (active initiator's) ACK to server (from **SYN_SENT** state) is lost?
2. Why the transition from **LISTEN** to **SYN_SENT**?
3. What if we get stuck in a state?
(*i.e.* what if the event that we are expecting never occurs?)
4. What if one side never sends a **FIN**?
5. Why is this **TIME_WAIT** state required?


TCP: The Time Wait State

When TCP receives ACK to last **FIN**, TCP enters **TIME_WAIT** state

- Protects future incarnations of connection from delayed segments
- **TIME_WAIT** = 2 x MSL
- Only valid segment that can arrive while in **TIME_WAIT** state is **FIN** retransmission
 - If such segment arrives, resend ACK & restart **TIME_WAIT** timer
- When timer expires, close TCP connection & delete connection record

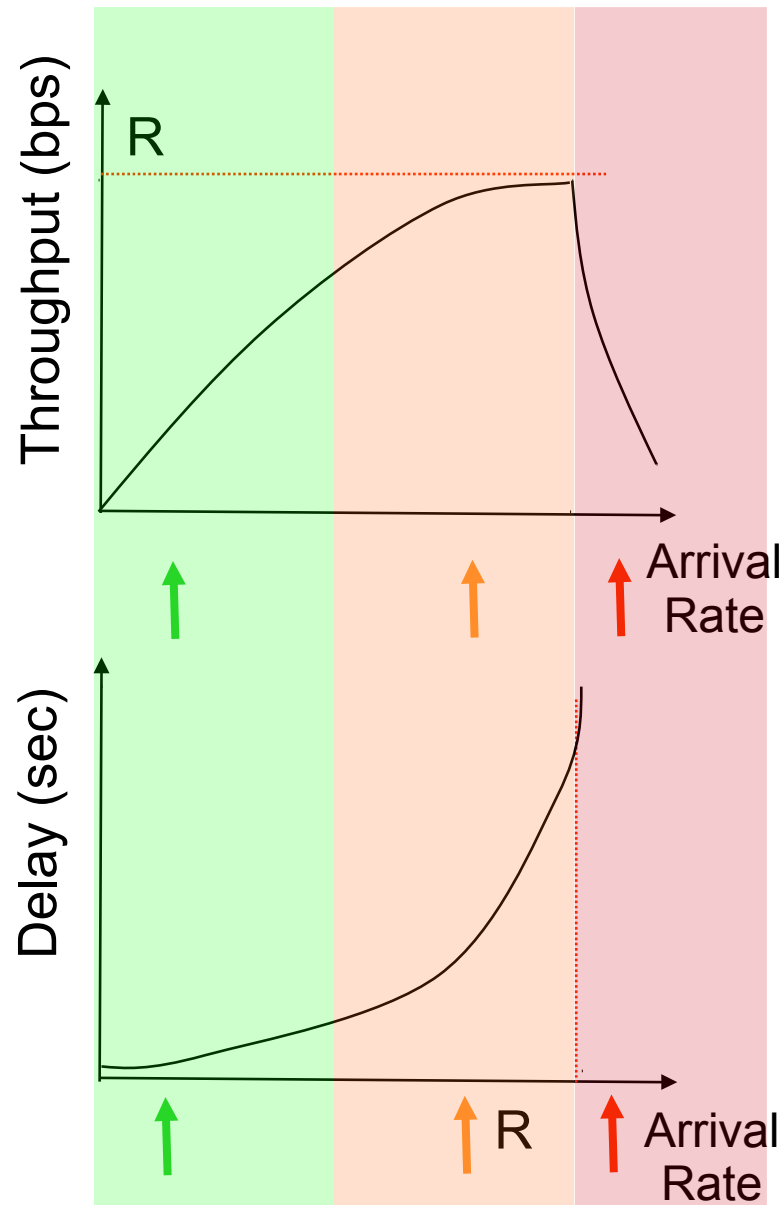
TCP: Congestion Control

- *Advertised window* size is used to prevent receiver's buffer overflow
- However, buffers at intermediate routers between source and destination may overflow ..**network congestion!**



- Congestion occurs when total arrival rate from all packet flows exceeds R over sustained period of time
- Buffers at multiplexer will fill up and packets will be dropped

Phases of Congestion



1. Light traffic

- Arrival Rate $\ll R$
- Low delay
- Can accommodate more

2. Knee (congestion onset)

- Arrival rate approaches R
- Delay increases rapidly
- Throughput begins to saturate

3. Congestive collapse

- Arrival rate $> R$
- Large delays, packet loss
- Throughput drops

Window Based Congestion Control

- Desired operating point: just before knee
 - Sources must control sending rates such that aggregate traffic avoids knee region
- Every TCP sender maintains *congestion window* **cwnd** to limit congestion at intermediate routers
- **Effective window*** is $\min(\text{cwnd}, \text{awnd}) - \text{data_in_flight}$
- Problem: source unaware of its “fair” share of available bandwidth
- Solution: adapt dynamically to available Bw
 - Sources probe the network by increasing **cwnd**
 - When congestion detected, sources reduce rate
 - Ideally, sources' sending rate stabilizes near desired range

* *Definition of effective window updated here*

TCP: Congestion Window Dynamics

TCP changes congestion window dynamically ..



- At **light traffic**: each segment is ACKed quickly
 - increase **cwnd** aggressively
- At **knee**: segment ACKs arrive, but more slowly
 - slow down increase in **cwnd**
- At **congestion**: (re)transmission timeouts occur due to delays and drops ..*sender gets duplicate ACKs*
 - reduce transmission rate, then probe again

TCP: Congestion Control - Slow Start

Slow start:

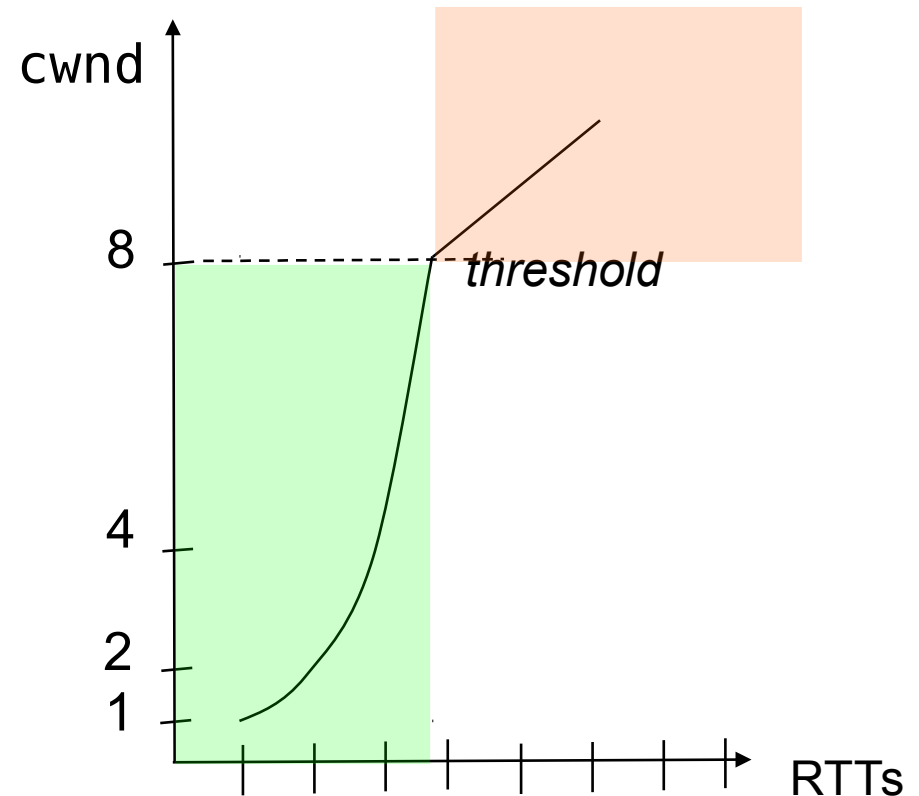
- (re)start tx with window size = 1 segment (MSS bytes)
- increase **cwnd** size by one segment each time ACK recvd
- congestion window increases exponentially !
- thus, not really slow at all!

Segments

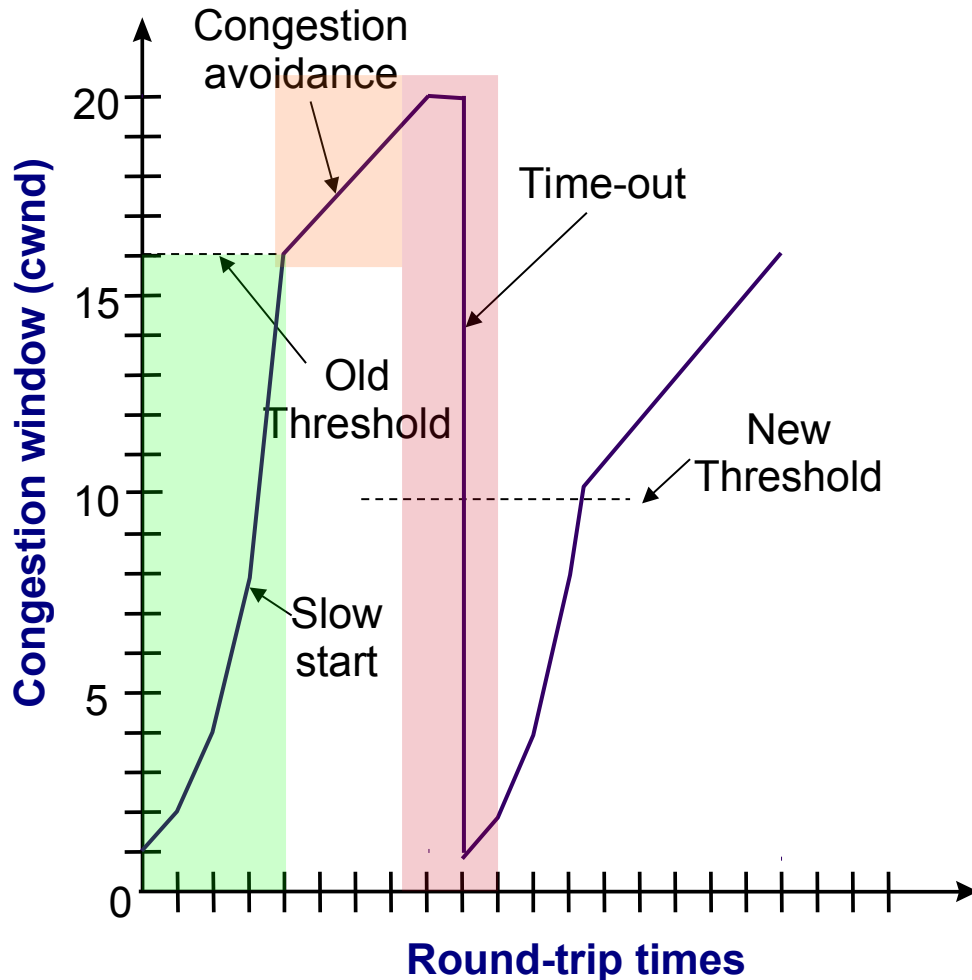


TCP: Congestion Control – Congestion Avoidance

- Algorithm progressively sets a *congestion threshold*
 - When $\text{cwnd} > \text{threshold}$, slow down rate of increase of cwnd
- Increase congestion window size by one segment per round-trip-time (RTT)
 - Each time an ACK arrives, cwnd increased by $1/\text{cwnd}$
 - In one RTT, cwnd segments are sent, so total increase in cwnd is $\text{cwnd} \times 1/\text{cwnd} = 1$
 - cwnd grows linearly with time



TCP: Congestion Control – Congestion Recovery



Timeout or duplicate ACKs means congestion in wired world

Eventually **cwnd** reaches bandwidth capacity. Then ..

On multiple timeouts or duplicate ACKs (many pkt drops in seq):


- Adjust congestion threshold* = $\frac{1}{2}$ x current **cwnd**
- Reset **cwnd** to 1
- Go back to slow-start

Over several cycles, threshold converges to about $\frac{1}{2}$ the max bandwidth

**During connection initialization, the threshold is set to an arbitrarily high value*


TCP: Fast Retransmit & Fast Recovery

- Congestion causes multiple segment drops
- If only single segment dropped, then subsequent segments trigger duplicate ACKs before timeout
- Can avoid large decrease in **cwnd** as follows:
 - When three duplicate ACKs arrive, retransmit lost segment immediately
 - Reset congestion threshold to $\frac{1}{2}$ **cwnd**
 - Reset **cwnd** to **congestion threshold + 3**
 - corresp. to the 3 duplicate ACKs
 - Remain in congestion avoidance phase
 - However if timeout occurs, reset **cwnd** to 1
 - In absence of timeouts, **cwnd** will oscillate around optimal value



Fast retransmit
Fast Recovery

TCP: Congestion Control - Fast Retransmit & Fast Recovery



TCP: Adaptive Retransmission

Issue: How to determine value of *retransmission timer*

Q: Why is this value critical?

A: Because,

- If set too high, lost pkts wait too long before being re-txed
... *this lowers line utilization & throughput*
- If set too low, pkts may time out and trigger re-tx unnecessarily
... *bandwidth wastage, low throughput*

Q: How best to determine RTT value(s)?

A: let us see..

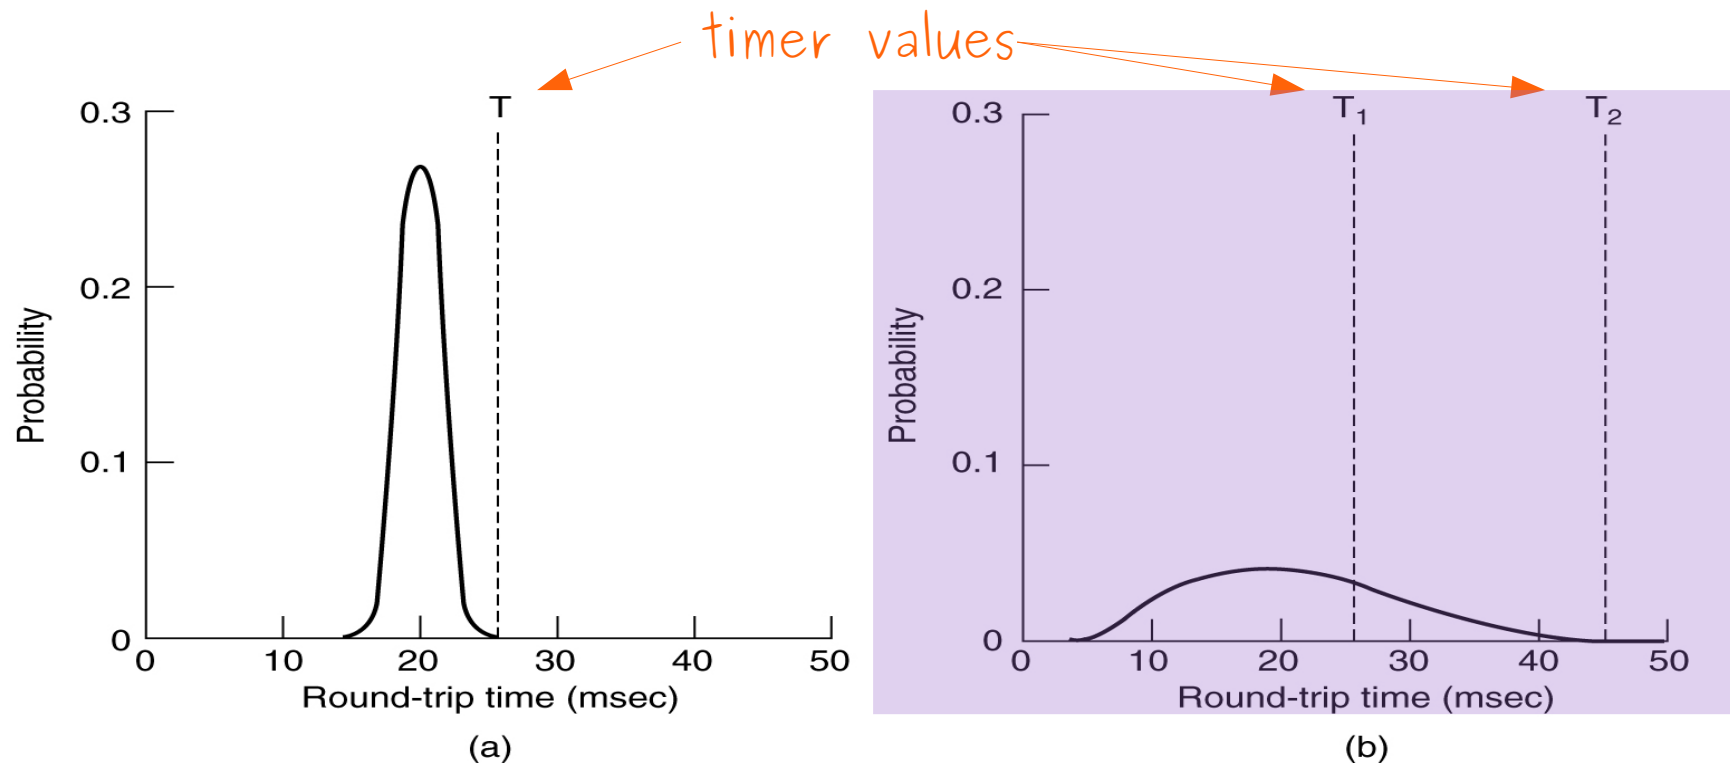
TCP: **Adaptive Retransmission** (contd.)

Q: How about simply calculating the mean of observed RTT values?

A: ?

TCP: Adaptive Retransmission (contd.)

What other pitfalls could we have in computing the retransmission timeout value?



(a) Probability Density of RTT (1-hop) as seen by the **data link layer**

(b) Probability Density of RTT (end-to-end) experienced by **TCP**
[The mean RTT value is 20 ms in both the above cases; and $T_1 = T$]

Notice the difficulty in setting the pkt timeout value for TCP

TCP: Adaptive Retransmission (contd.)

(RFC 793 - Exponential Averaging)

Estimated Round-Trip Time (ERTT):

$$\text{ERTT}(K+1) = \alpha \times \text{ERTT}(K) + (1 - \alpha) \times \text{SampleRTT}(K)$$

where $0 < \alpha < 1$ (suggested value is $0.8 < \alpha < 0.9$)


Idea

The older the observation, the lower its weight in the average...
i.e. “*exponential forgetting*”

And, $\text{RTO} = 2 \times \text{ERTT}$

(RTO = Retransmission TimeOut)

TCP: Adaptive Retransmission (contd.)



A problem of association!

– why **SampleRTT** computation is not trivial

TCP: Adaptive Retransmission (contd.)

Stated simply ...

- If ACK received for retransmitted segment, then 2 possibilities:
 - ACK is for first tx, *or*
 - ACK is for second tx

TCP source cannot distinguish betwn above 2 cases

To sidestep this problem, we use the **Karn-Partridge** algorithm (*see next slide*)

TCP: Adaptive Retransmission (contd.)

(Karn-Partridge Algorithm)

- On segment re-tx, stop measuring SampleRTT
- Increase RTO on each segment retransmission (i.e. **backoff** process)

$$RTO_{i+1} = q \times RTO_i, \text{ where } q = \text{constant}$$

if $q = 2$, it is called **binary exponential backoff**

TCP: Adaptive Retransmission (contd.)

(Jacobson/Karels Algorithm)

$$\text{Difference}(K) = \text{SampleRTT}(K) - \text{ERTT}(K)$$

$$\text{ERTT}(K + 1) = \text{ERTT}(K) + [\delta \times \text{Difference}(K)]$$

$$\text{Deviation}(K + 1) = \text{Deviation}(K) + \delta \times [|\text{Difference}(K)| - \text{Deviation}(K)]$$

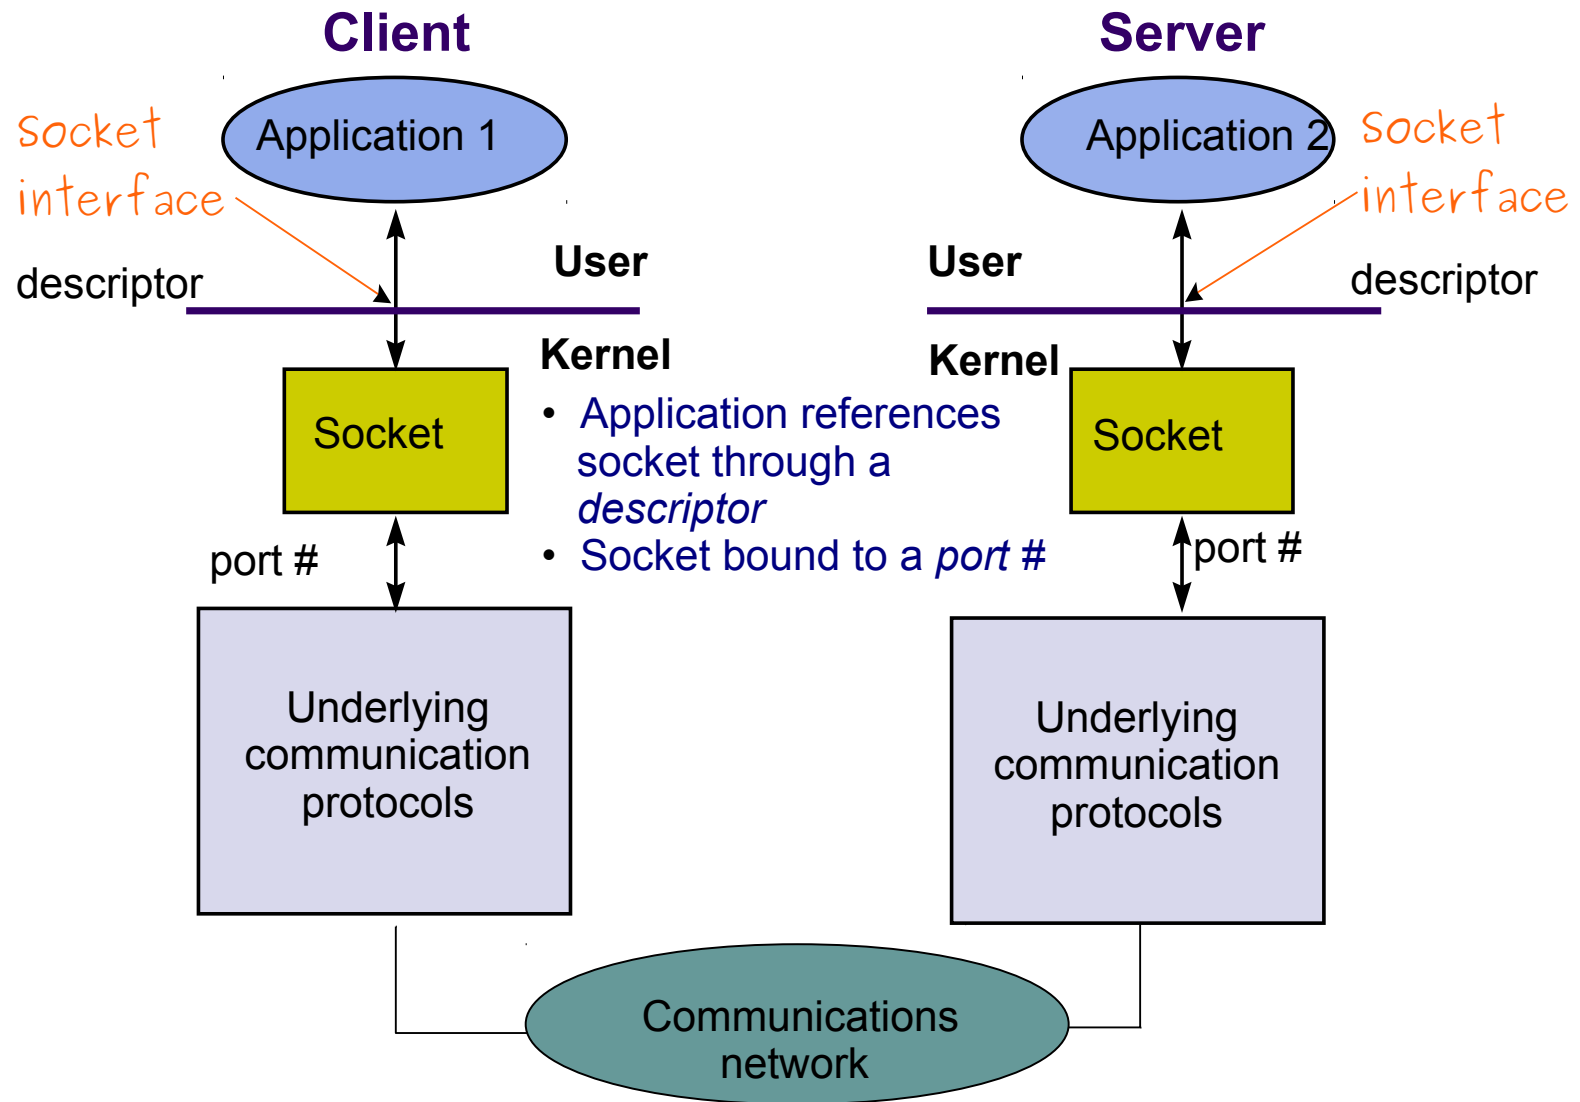
$$\text{RTO}(K + 1) = \mu \times \text{ERTT}(K+1) + \Phi \times \text{Deviation}(K + 1)$$

where: $0 < \delta < 1$, $\mu = 1$ (usually), and $\Phi = 2$ or 4
(most current implementations use $\Phi = 4$)

Berkeley Sockets: The Socket API

- API (Application Programming Interface)
 - Provides std functions that can be called by applications
- Berkeley UNIX Sockets API
 - Abstraction for applications to send or receive data; hides details of underlying protocols & mechanisms
 - Applications create sockets that “plug into” network
 - Applications write/read to/from sockets
 - Implemented in OS kernel
- Also in Windows (WinSock) , Linux, and other OSes

Communications through Socket Interface



Stream Mode of Service

Connection-oriented (w/ TCP)

- First, **set up** connection between two peer application processes
- Then, reliable bidirectional in-sequence transfer of *byte stream* (boundaries not preserved in transfer)
- Multiple write/read between peer processes
- Finally, **connection release**

Connectionless (w/ UDP)

- Immediate transfer of one block of information (boundaries preserved)
- No setup overhead & delay
- Destination address with each block
- Send/receive to/from multiple peer processes
- Possibly out-of-order, possible pkt loss

Client & Server Differences

- **Server**

- Specifies well-known port # when creating socket
- May have multiple IP addresses (net interfaces)
- Waits passively for client requests

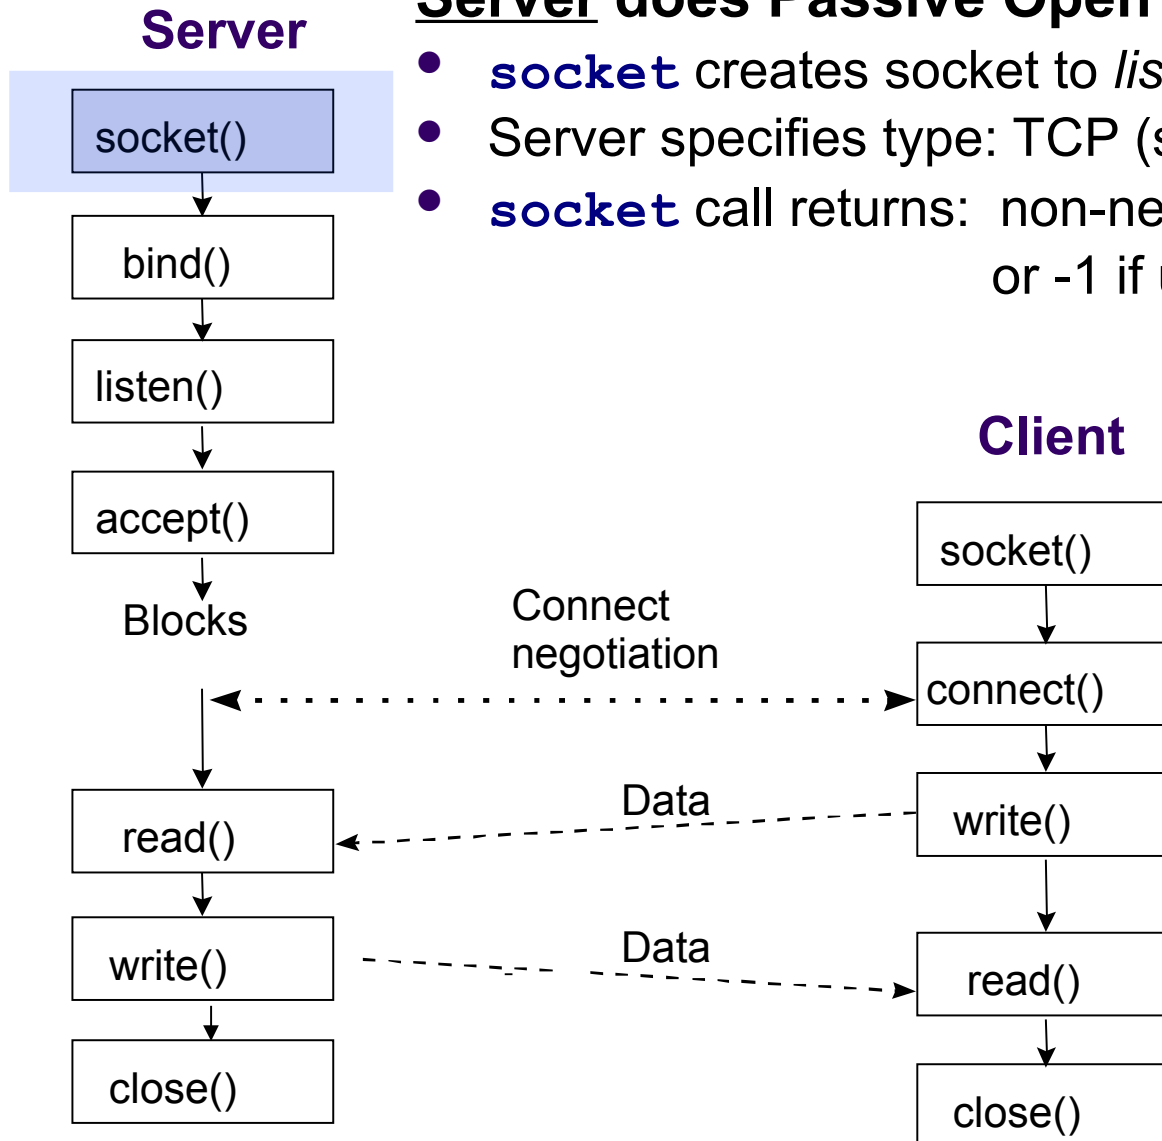
- **Client**

- Assigned ephemeral port #
- Initiates communications with server
- Needs to know server's IP address & port #
 - or, (logical name + DNS support) and (protocol)
- Server learns client's address & port # from request pkt

Socket Calls for Connection-Oriented Mode: `socket()` page 1

Server does Passive Open


- `socket` creates socket to *listen* for connection requests
- Server specifies type: TCP (stream)
- `socket` call returns: non-negative integer *descriptor*; or -1 if unsuccessful



Socket Call: `bind()`

page 2


Server



Server does Passive Open

- **bind** assigns local address & port # to socket with specified descriptor
- Wildcard IP address for multiple net interfaces
- **bind** call returns: 0 (success); or -1 (failure)
- Failure if port # already in use and reuse option not set

Client



Connect negotiation


Data

Data

Socket Call: `listen()`

page 3


Server



Server does Passive Open

- `listen` indicates to TCP readiness to receive connection requests for socket with given descriptor
- Parameter specifies max number of requests that may be queued while waiting for server to accept them
- `listen` call returns: 0 (success); or -1 (failure)

Client



Connect negotiation


Data

Data

Socket Call: `accept()`

page 4


Server



Server does Passive Open

- Server calls **accept** to accept incoming requests
- **accept** blocks if queue is empty

Client




Connect negotiation

Data

Data

Socket Call: Client's Side `socket()` page 5

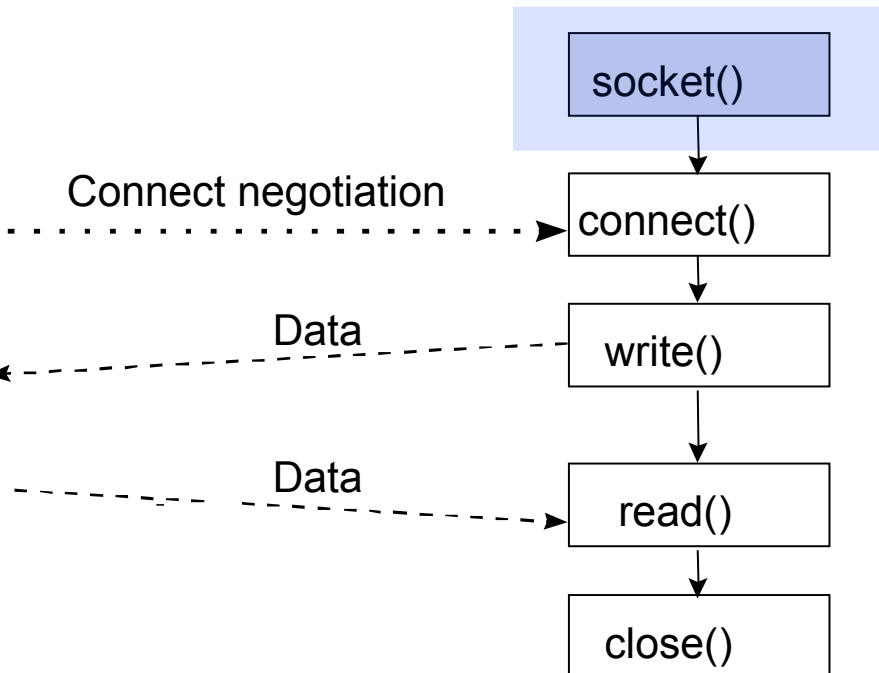
Server



Client does Active Open


- `socket` creates socket to connect to server
- Client specifies type: TCP (stream)
- `socket` call returns: non-negative integer *descriptor*; or -1 if unsuccessful

Client



Socket Call: connect() page 6


Server



Client does Active Open

- **connect** establishes a connection on the local socket (with specified descriptor) to the specified remote address and port #
- **connect** returns 0 if successful; -1 if unsuccessful

Client



Connect negotiation

Data

Data

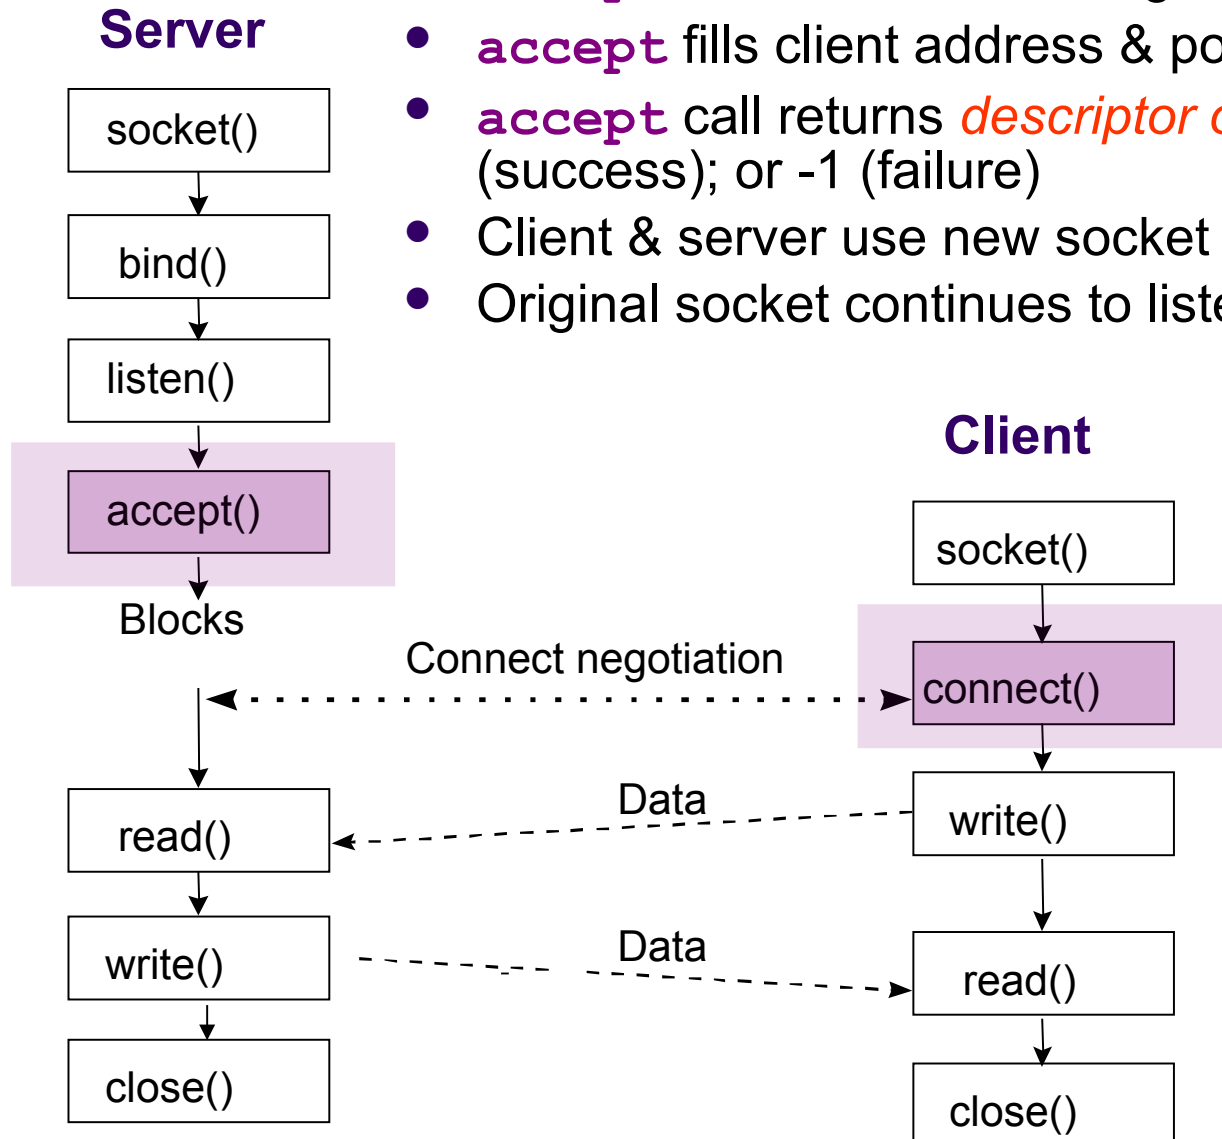
Note:

connect initiates TCP's three-way handshake

Socket Call: `accept()`

page 7

- `accept` wakes with incoming connection request
- `accept` fills client address & port # into address structure
- `accept` call returns *descriptor of new connection socket* (success); or -1 (failure)
- Client & server use new socket for data transfer
- Original socket continues to listen for new requests

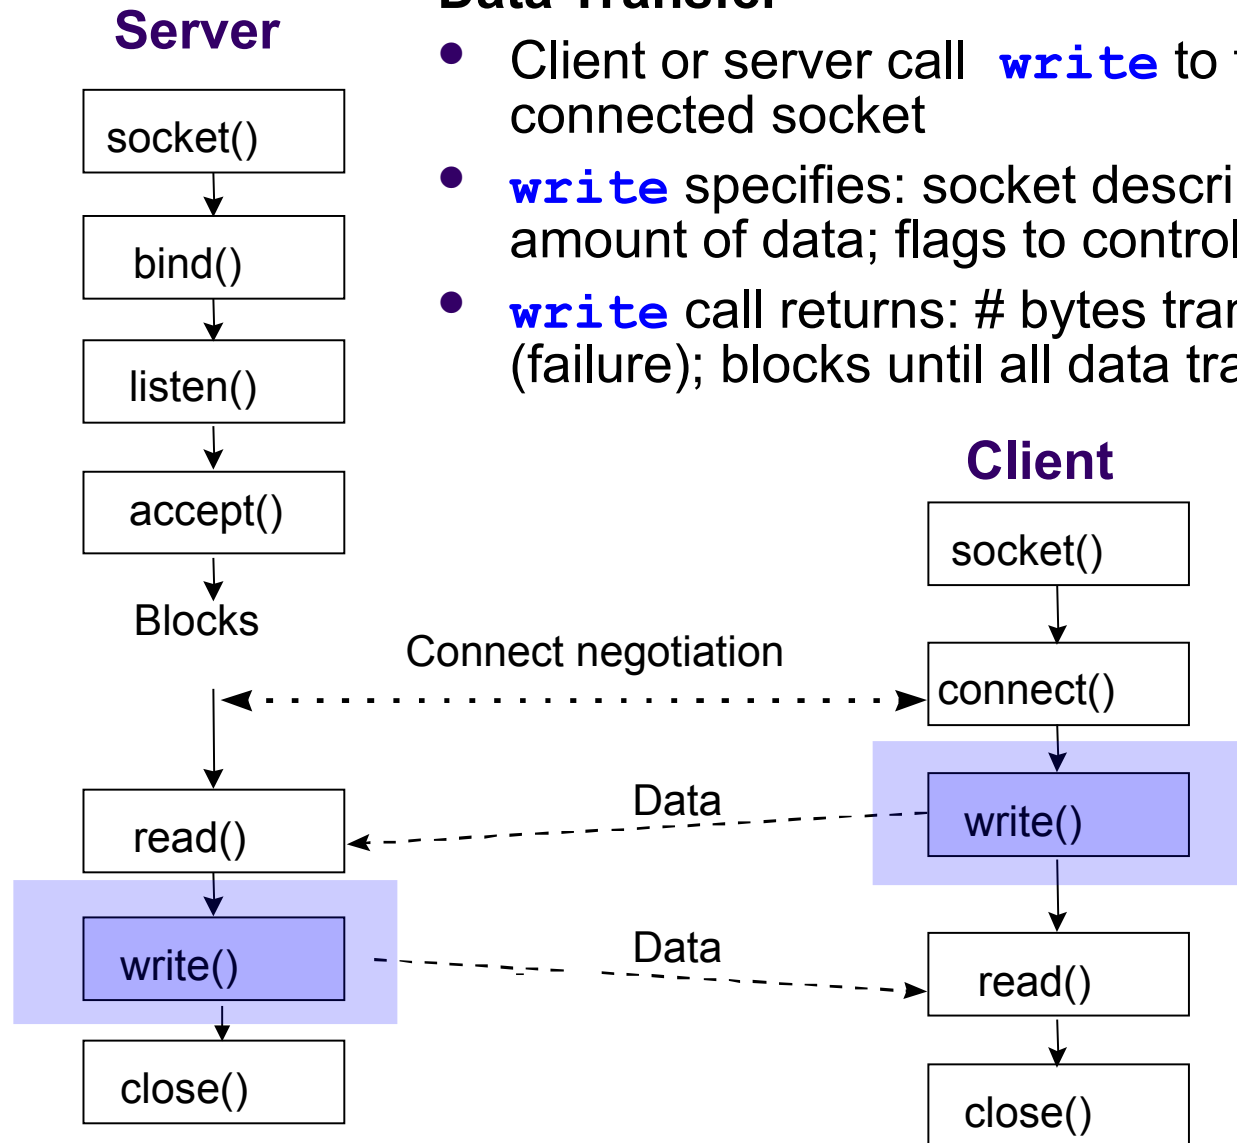


Socket Call: `write()`

page 8

Data Transfer

- Client or server call `write` to transmit data into a connected socket
- `write` specifies: socket descriptor; pointer to a buffer; amount of data; flags to control transmission behavior
- `write` call returns: # bytes transferred (success); or -1 (failure); blocks until all data transferred

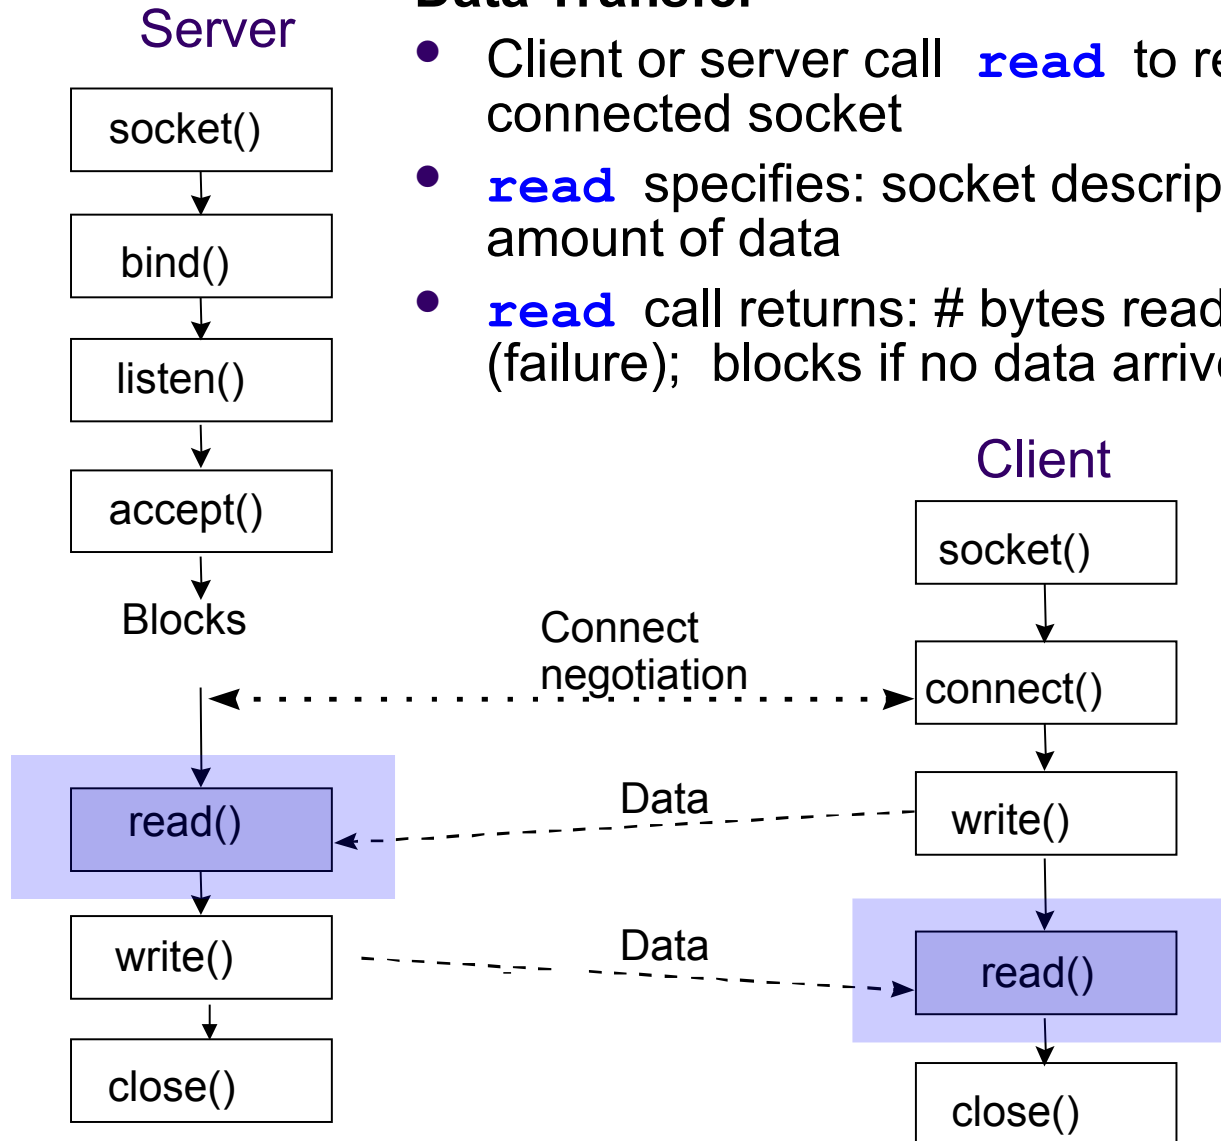


Socket Call: `read()`

page 9

Data Transfer

- Client or server call `read` to receive data from a connected socket
- `read` specifies: socket descriptor; pointer to a buffer; amount of data
- `read` call returns: # bytes read (success); or -1 (failure); blocks if no data arrives




Note:

`write` and `read` can be called multiple times to transfer byte streams in both directions

Socket Call: `close()`

page 9


Server



Connection Termination

- Client or server call **close** when socket is no longer needed
- **close** specifies the socket descriptor
- **close** call returns: 0 (success); or -1 (failure)

Client



Note:

close initiates TCP's graceful close sequence

Example: TCP Echo Server *home reading from this point on*

```
/* A simple echo server using TCP */
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define SERVER_TCP_PORT 3000
#define BUFLLEN 256

int main(int argc, char **argv)
{
    int n, bytes_to_read;
    int sd, new_sd, client_len, port;
    struct sockaddr_in server, client;
    char *bp, buf[BUFLLEN];

    switch(argc) {
        case 1:
            port = SERVER_TCP_PORT;
            break;
        case 2:
            port = atoi(argv[1]);
            break;
        default:
            fprintf(stderr, "Usage: %s [port]\n",
                argv[0]);
            exit(1);
    }

    /* Create a stream socket */
    if ((sd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
        fprintf(stderr, "Can't create a socket\n");
        exit(1);
    }
}
```

```
/* Bind an address to the socket */
bzero((char *)&server, sizeof(struct sockaddr_in));
server.sin_family = AF_INET;
server.sin_port = htons(port);
server.sin_addr.s_addr = htonl(INADDR_ANY);
if (bind(sd, (struct sockaddr *)&server,
    sizeof(server)) == -1) {
    fprintf(stderr, "Can't bind name to socket\n");
    exit(1);
}

/* queue up to 5 connect requests */
listen(sd, 5);

while(1) {
    client_len = sizeof(client);
    if ((new_sd = accept(sd, (struct sockaddr *)&client,
        &client_len)) == -1) {
        fprintf(stderr, "Can't accept client\n");
        exit(1);
    }

    bp = buf;
    bytes_to_read = BUFLLEN;
    while ((n = read(new_sd, bp, bytes_to_read)) > 0) {
        bp += n;
        bytes_to_read -= n;
    }
    printf("Rec'd: %s\n", buf);

    write(new_sd, buf, BUFLLEN);
    printf("Sent: %s\n", buf);
    close(new_sd);
}
close(sd);
return(0);
} /*end*/
```

Example: TCP Echo Client

```
/* A simple TCP client */
#include <stdio.h>
#include <netdb.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define SERVER_TCP_PORT 3000
#define BUFLLEN 256

int main(int argc, char **argv)
{
    int n, bytes_to_read;
    int sd, port;
    struct hostent *hp;
    struct sockaddr_in server;
    char *host, *bp, rbuf[BUFLLEN], sbuf[BUFLLEN];

    switch(argc) {
    case 2:
        host = argv[1];
        port = SERVER_TCP_PORT;
        break;
    case 3:
        host = argv[1];
        port = atoi(argv[2]);
        break;
    default:
        fprintf(stderr, "Usage: %s host [port]\n",
            argv[0]);
        exit(1);
    }

    /* Create a stream socket */
    if ((sd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
        fprintf(stderr, "Can't create socket\n");
        exit(1);
    }

    bzero((char *)&server, sizeof(struct sockaddr_in));
    server.sin_family = AF_INET;
    server.sin_port = htons(port);
    if ((hp = gethostbyname(host)) == NULL) {
        fprintf(stderr, "Can't get server's address\n");
        exit(1);
    }
    bcopy(hp->h_addr, (char *)&server.sin_addr, hp->h_length);

    /* Connecting to the server */
    if (connect(sd, (struct sockaddr *)&server,
        sizeof(server)) == -1) {
        fprintf(stderr, "Can't connect\n");
        exit(1);
    }
    printf("Connected: server's address is %s\n", hp->h_name);

    printf("Transmit:\n");
    gets(sbuf);
    write(sd, sbuf, BUFLLEN);

    printf("Receive:\n");
    bp = rbuf;
    bytes_to_read = BUFLLEN;
    while ((n = read(sd, bp, bytes_to_read)) > 0) {
        bp += n;
        bytes_to_read -= n;
    }
    printf("%s\n", rbuf);

    close(sd);
    return(0);
}
```

Socket Calls for Connection-Less Mode (UDP)

Read this on your own from Ch 2 of the textbook

